Amogh **Akshintala**

Ph.D. Student, Dept. of Computer Science, The University of North Carolina at Chapel Hill

6906 Guadalupe Street, #304, Austin, TX 78752

+1 (631)-371-4820 | ✉ aakshintala@cs.unc.edu | 🏠 aakshintala.com | 🐙 aakshintala

Education

University of North Carolina - Chapel Hill

Doctor of Philosophy   *Aug 2016 — Ongoing*

Stony Brook University

Doctor of Philosophy (Transferred to UNC)   *Jan 2014 — Aug 2016*

M.S. in Computer Science (awarded May 2016)   *Aug 2012 — Dec 2013*

Visvesvaraya Technological University

B.E. in Computer Science and Engineering   *2008 — 2012*

Publications

Conferences

Trillium: The code is the IR

Amogh Akshintala, Hangchen Yu, Arthur Peters, Christopher J. Rossbach   *VIRT '19*

x86-64 Instruction Usage among C/C++ Applications

Amogh Akshintala, Bhushan P. Jain, Chia-che Tsai, Michael Ferdman, Donald E. Porter   *SYSTOR '19*

Optimizing Every Operation in a Write-optimized File System   *Awarded Best Paper*

Jun Yuan, Yang Zhan, William Jannen, Prashant Pandey, Amogh Akshintala, Kanchan Chandnani,
Pooja Deo, Zardosht Kasheff, Leif Walsh, Michael Bender, Martin Farach-Colton, Rob Johnson,   *FAST '16*
Bradley C. Kuszmaul, and Donald E. Porter.

BetrFS: A Right-Optimized Write-Optimized File System   *Best Paper Runner-up*

William Jannen, Jun Yuan, Yang Zhan, Amogh Akshintala, John Esmet, Yizheng Jiao, Ankur Mittal,
Prashant Pandey, Phaneendra Reddy, Leif Walsh, Michael Bender, Martin Farach-Colton, Rob   *FAST '15*
Johnson, Bradley C. Kuszmaul, and Donald E. Porter.

Journals

Writes Wrought Right, and Other Adventures in File System Optimization.

Jun Yuan, Yang Zhan, William Jannen, Prashant Pandey, Amogh Akshintala, Kanchan Chandnani,
Pooja Deo, Zardosht Kasheff, Leif Walsh, Michael A. Bender, Martin Farach-Colton, Rob Johnson,   *ACM Transactions on Storage (TOS)*
Bradley C. Kuszmaul, and Donald E. Porter.   *Mar'17*

BetrFS: Write-Optimization in a Kernel File System

William Jannen, Jun Yuan, Yang Zhan, Amogh Akshintala, John Esmet, Yizheng Jiao, Ankur Mittal,
Prashant Pandey, Phaneendra Reddy, Leif Walsh, Michael Bender, Martin Farach-Colton, Rob   *ACM Transactions on Storage (TOS)*
Johnson, Bradley C. Kuszmaul, and Donald E. Porter.   *Nov'15*

Workshops

**USETL: Unikernels for Serverless Extract Transform and Load. Why should you
settle for less?**   *Awarded Best Paper*

Henrique Fingler, Amogh Akshintala, Christopher J. Rossbach   *APSys '19*

Automatic Virtualization of Accelerators

Hangchen Yu, Arthur Peters, Amogh Akshintala, and Christopher J. Rossbach.   *HotOS '19*

**Talk to My Neighbors Transport: Decentralized Data Transfer and Scheduling
Among Accelerators.**

Amogh Akshintala, Vance Miller , Donald E. Porter, and Christopher J. Rossbach.   *SFMA '18*

Experience

Internships

VMware Research Group   *Palo Alto, CA.*

Research Intern   *May 2018 — Aug. 2018*

VMware Research Group   *Austin, TX.*

Research Intern   *May 2017 — Aug. 2017*

VMware Inc.   *Cambridge, MA.*

Member of Technical Staff - Intern   *May 2014 — Aug. 2014*

Tintri Inc.   *Mountain View, CA*

Member of Technical Staff - Intern   *May 2013 - Aug. 2013*

Academic

University of Texas at Austin

Visiting Researcher   *Jan 2019 — Dec 2019*

University of Texas at Austin

Visiting Researcher   *Fall 2017*

UNC Department of Computer Science

Research Assistant   *Aug 2016 — present*

Stony Brook University

Research Assistant   *Jan 2015 — Aug. 2016*

Stony Brook University   *Operating Sytems (Graduate)*

Teaching Assistant   *Fall 2014*

Stony Brook University   *Network Security (Graduate)*

Teaching Assistant   *Spring 2014*

Service

External Reviewer

ACM SOCC '18

Projects

TMNT: accelerating data movement among accelerators   *Spring '18 - ongoing*

- Data movement is a first-order determinant of performance when programming high-throughput accelerators. When you throw in the additional challenge of co-ordination among multiple accelerators, the problem is compounded because of the synchronous nature of control in the Master-Slave model that most accelerators operate under.
- We're designing an capability-based hardware structure that provides the necessary primitives to express data placement/movement, manage co-ordination and scheduling of computation on accelerators, and enforce capabilities for processes running on accelerators. We believe these extensions naturally fit the data-flow programming paradigm, which should greatly ease programmability.
- Outcome: 1 Workshop paper. Work ongoing.

Accelerator Virtualization   *Sprint '18 - ongoing*

- Accelerators seem to operate as horizontally and vertically isolated silos — i.e., the only exposed surfaces that can be interposed on are either the user-space API or the HW interface. Leveraging this key insight, we're building an accelerator-virtualization framework that automatically generates much of the code required to forward user-space accelerator APIs through the hypervisor in order to provide the desirable properties of virtualization.
- Outcome: 1 workshop paper. 1 paper under submission. Work ongoing.

GPGPU Virtualization *Summer '17 - Spring '18*

- Many methods have been proposed to virtualize general purpose compute on GPUs. However, none of them hit the right spot. We observe that there are actually two separate elements that must be virtualized when dealing with GPGPUs: device control, and compute. We extend the para-virtual model used by VMware to investigate our hypothesis that handling these two elements separately is the key to achieving good GPGPU performance.
- Outcome: 1 conference paper

Instruction Popularity *Summer '16 — ongoing*

- Overlapping-ISA multi-core computers have been actively studied in the past decade. However, in most of these studies the data, used to craft the various subsets of the ISA, is collected from a small number of applications, usually popular, and often outdated, benchmarks.
- We performed static analysis on 9000 C/C++ binaries in the Ubuntu 16.04 repositories to get a higher fidelity understanding of static instruction distribution among applications.
- Outcome: 1 Conference paper. Visualization tool: x86instructionpop.com

BetrFS (http://www.betrfs.org/) *Spring '14 — Spring '16*

- The B^ϵ–tree File System, or BetrFS, is an in-kernel file system that uses B^ϵ–trees to organize on-disk storage. B^ϵ–trees are a form of write-optimized dictionaries, and offer the same asymptotic behavior for sequential I/O and point queries as a B-tree. The advantage of a B^ϵ–tree is that it can also ingest small, random writes 1-2 orders of magnitude faster than B-trees and other standard on-disk data structures. The goal of BetrFS is to realize performance that strictly dominates the performance of current, general-purpose file systems.
- I was the Benchmark Czar for the project — I was responsible for measuring and understanding performance of BetrFS.
- Outcome: 2 Conference and 2 Journal papers. Project still ongoing; my involvement is limited.

Investigating Next Page Prefetcher behaviour for virtual machines *Spring 2014 - Summer 2015*

- Significant differences in the impact of prefetching on the performance of applications running in a VM (vs native) triggered an investigation of the behaviour of the Next Page Prefetcher, which had just been incorporated into Intel CPUs.
- Outcome: We were unable to discern the source of the variation, but I learnt a lot about caches, prefetching, and profiling tools.

Reduced x86_64 backend for LLVM compiler *Spring 2013 - Spring 2015*

- As a first step towards the exploring the design-space of ISA-homogeneous, Capability-heterogenous multi-core processors, we attempted to build a backend for the LLVM Compiler that operated on a configurable subset of x86_64 ISA.
- I mentored 2 M.S. students working on this project.