

# **COMP 590-154:**

# **Computer Architecture**

Speculation and Traps in Out-of-Order Cores

# What is wrong with Tomasulo's?

- Branch instructions
  - Need *branch prediction* to guess what to fetch next
  - Need *speculative execution* to “clean up” wrong guesses
- Exceptions and Traps (“software” interrupts)
  - Need to handle uncommon execution cases
    - Jump to a software handler
  - Should follow the insn. on which they were triggered
  - Often referred to as *precise interrupts*

Don't know relative order of instructions in RS

# Speculation and Precise Interrupts

- When branch is mis-speculated by predictor
  - Must reset state (e.g., regs) to time of branch
- Sequential semantics for interrupts
  - All insns. before interrupt should be complete
  - All insns. after interrupt should look as if never started (abort)
- What makes this difficult?
  - Younger insns. finish before branch → must undo writebacks
  - Older insns. not done when young branch resolves → must wait
    - Older insn. takes page fault or divide by zero → forget the branch

Same problem → Same solution

# Precise State

- Speculative execution requires
  - (Ability to) abort & restart at every branch
  - Abort & restart at every load (covered in later lecture)
- Synchronous (exception and trap) events require
  - Abort & restart at every load, store, divide, ...
- Asynchronous (hardware) interrupts require
  - Abort & restart at every ??
- Real world: bite the bullet
  - Implement abort & restart at every insn.
  - Called *precise state*

# Precise State Implementation Options

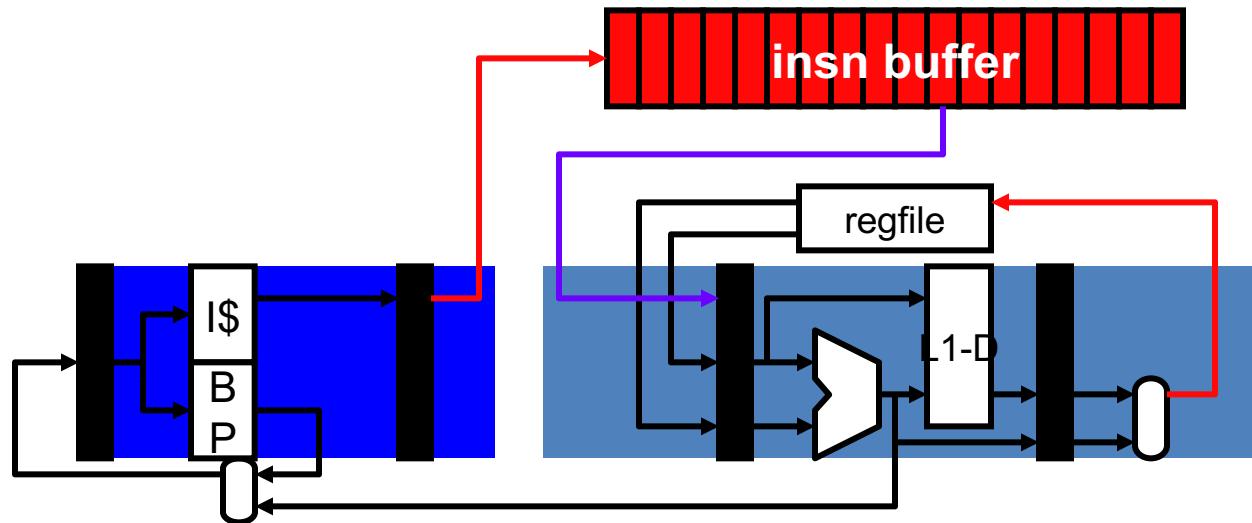
- Imprecise state: ignore the problem!
  - Makes page faults (any restartable exceptions) difficult
  - Makes speculative execution practically impossible
- Force in-order completion (W): stall pipe if necessary
  - Slow (takes away benefit of Out-of-Order)
- Keep track of precise state in hardware
  - Reset current state from precise state when needed

Everything is better in hardware

# Our-of-Order Topics

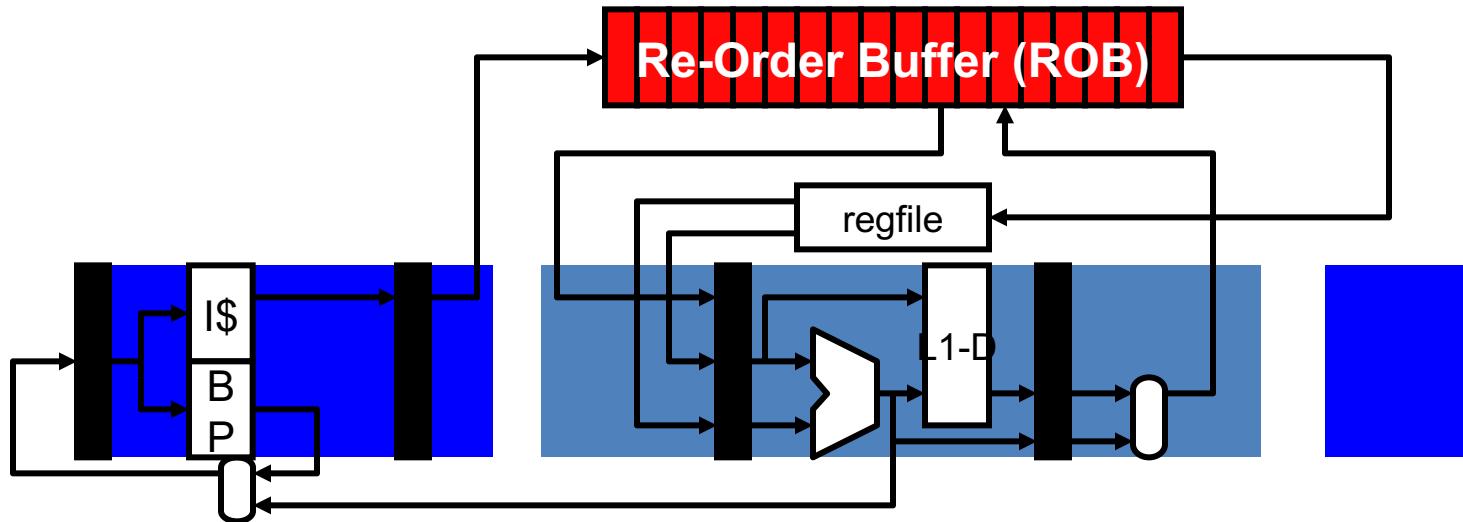
- “Scoreboarding”
  - First OoO, no register renaming
- “Tomasulo’s algorithm”
  - OoO with register renaming
- Handling precise state and speculation
  - P6-style execution (Intel Pentium Pro)
  - R10k-style execution (MIPS R10k)
- Handling memory dependencies

# The Problem with Precise State



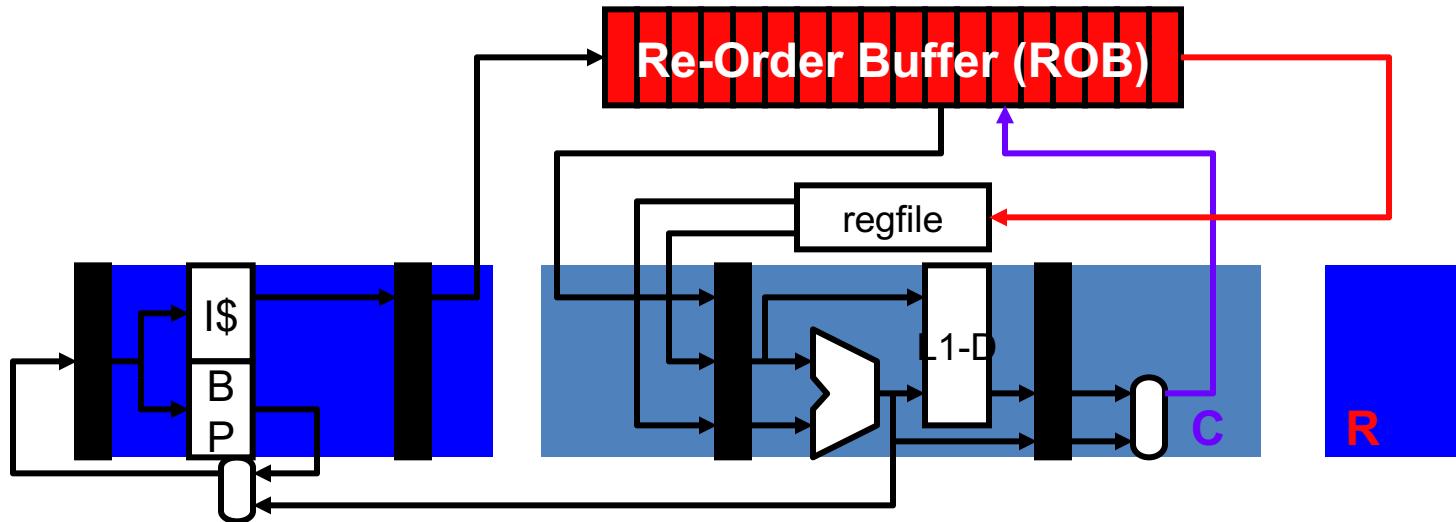
- Problem: writeback combines two functions
  - Forward values to younger insns.: out-of-order is OK
  - Write values to registers: needs to be in order
- Similar solution as for OoO decode
  - Split writeback into two stages

# Re-Order Buffer (ROB)



- Insn. buffer → Re-Order Buffer (ROB)
  - Buffer completed results en route to register file
  - Can be merged with RS (RUU) or separate (common today)
- Split writeback (W) into two stages
  - Why is there no latch between W1 and W2?

# Complete and Retire

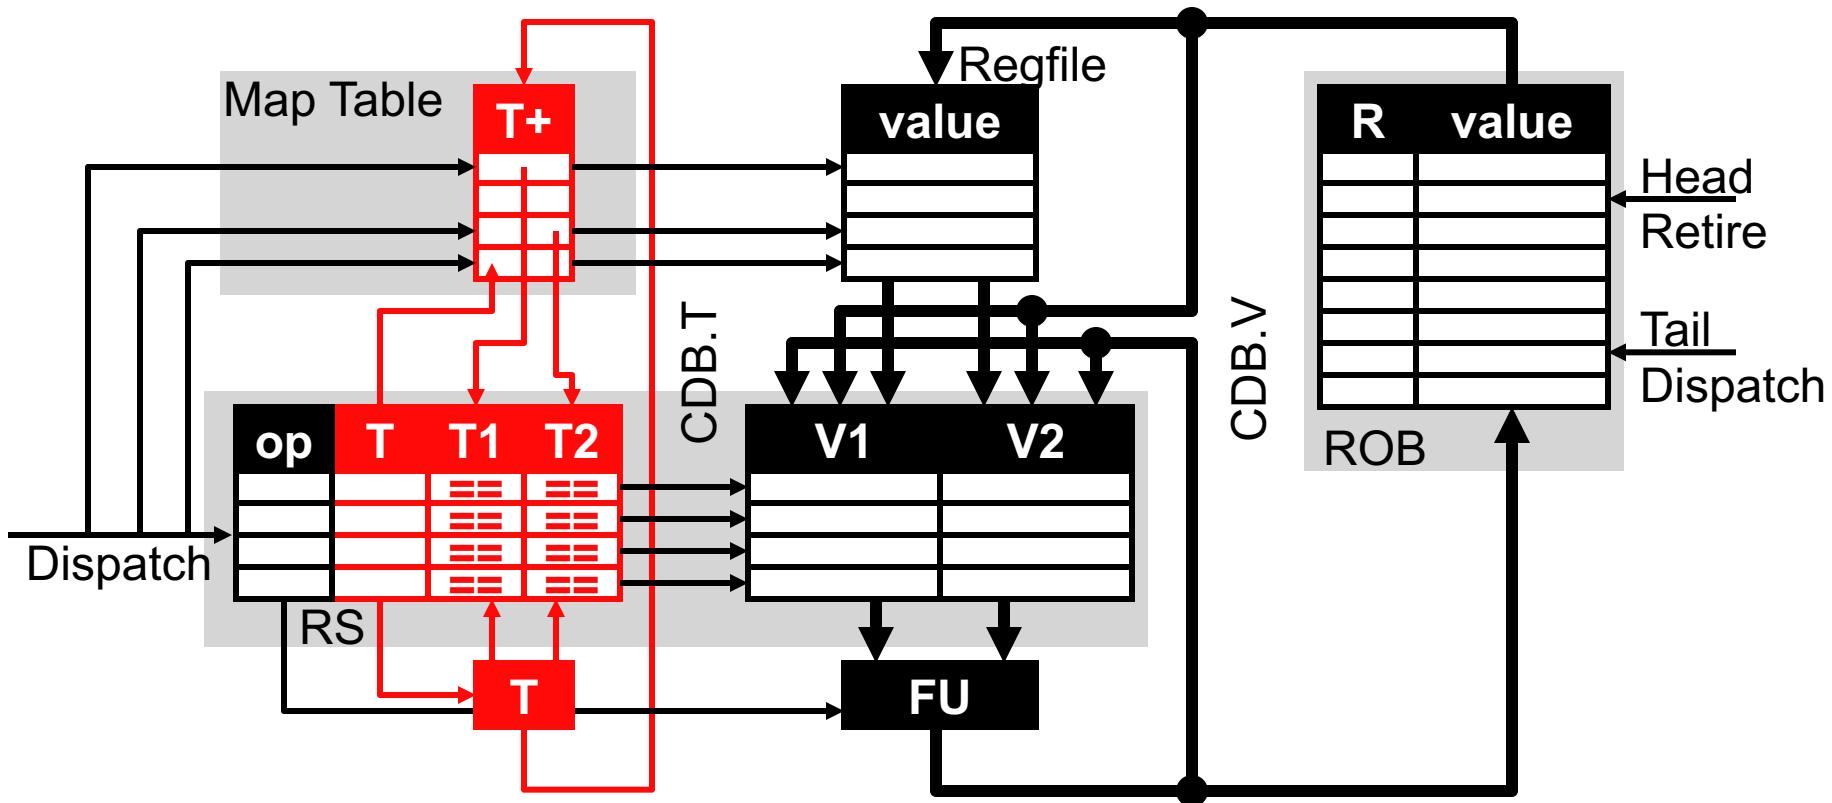


- Complete (**C**): insns. write results into ROB
  - Out-of-order: don't block younger insns.
- Retire (**R**): a.k.a. commit
  - ROB writes results to register file
  - In order: stall back-propagates to younger insns.

# P6 Data Structures

- P6: Start with Tomasulo's algorithm... add ROB
- ROB (separate from RS)
  - **head, tail**: pointers maintain sequential order
  - **R**: insn. output register, **V**: insn. output value
- Tags are different
  - Tomasulo: RS# → P6: ROB#
- Map Table is different
  - **T+**: tag + “ready-in-ROB” bit
    - $T==0 \rightarrow$  Value is ready in register file
    - $T!=0 \rightarrow$  Value is not ready
    - $T!=0+ \rightarrow$  Value is ready in the ROB

# P6 Data Structures (1/2)



# P6 Data Structures (2/2)

ROB							Map Table		CDB		
ht	#	Insn	R	V	S	X	C	Reg	T+	T	V
	1	<code>ldf X(r1), f1</code>						<code>f0</code>			
	2	<code>mulf f0, f1, f2</code>						<code>f1</code>			
	3	<code>stf f2, Z(r1)</code>						<code>f2</code>			
	4	<code>addi r1, 4, r1</code>						<code>r1</code>			
	5	<code>ldf X(r1), f1</code>									
	6	<code>mulf f0, f1, f2</code>									
	7	<code>stf f2, Z(r1)</code>									

Reservation Stations									
#	FU	busy	op	T	T1	T2	V1	V2	
1	ALU	no							
2	LD	no							
3	ST	no							
4	FP1	no							
5	FP2	no							

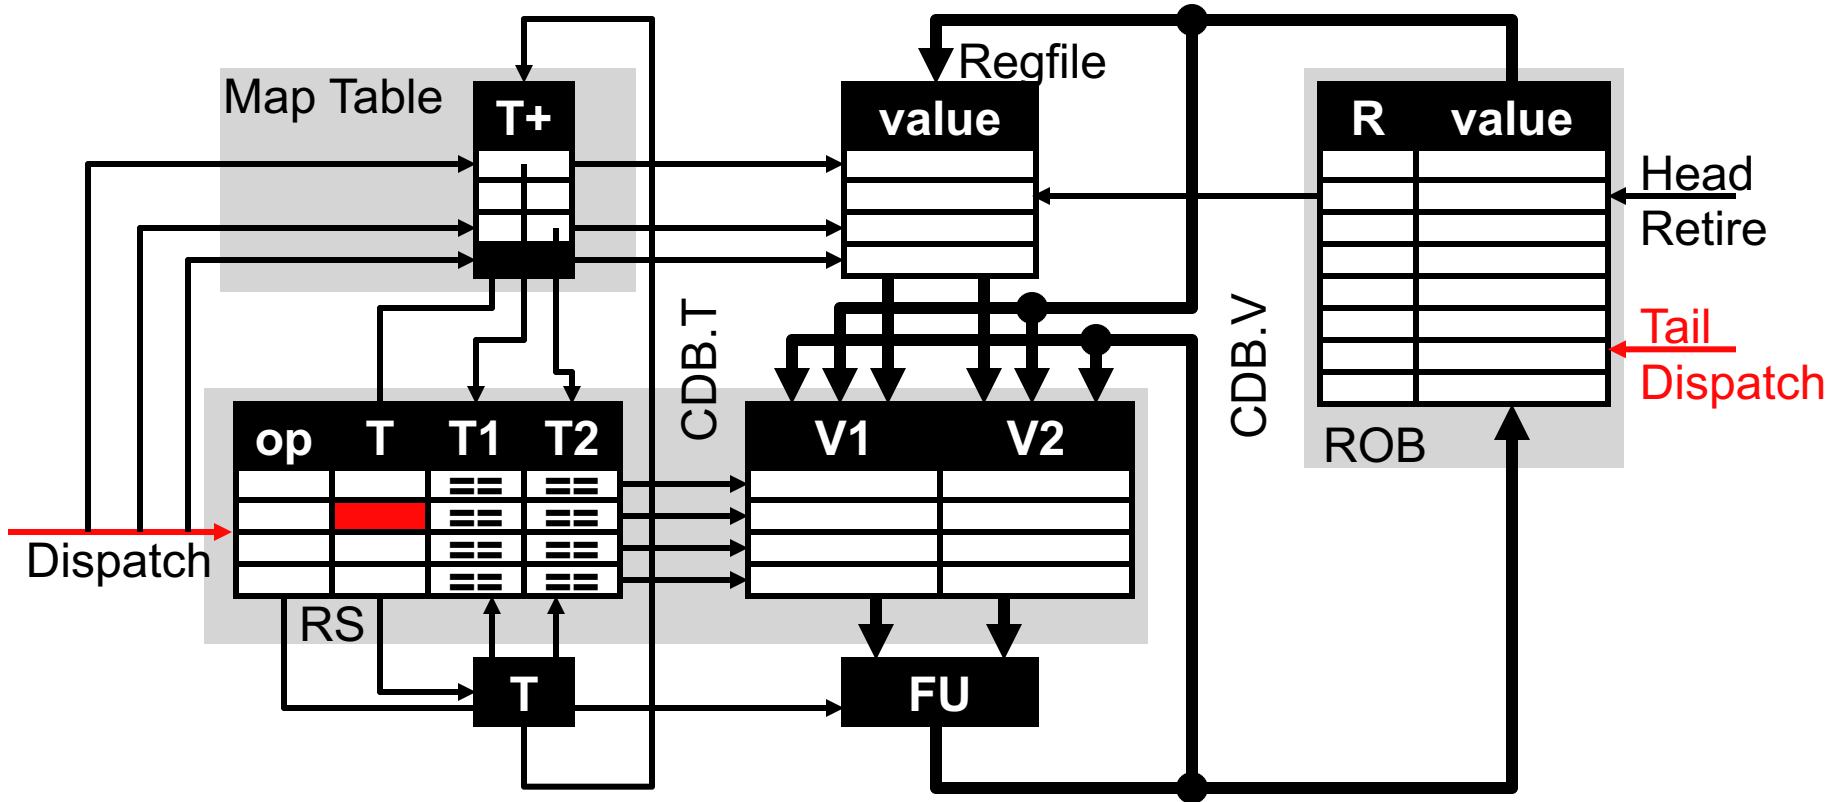
# P6 Pipeline

- New pipeline structure: F, **D**, S, **X**, **C**, **R**
  - **D (dispatch)**
    - Structural hazard (ROB/RS) ? **stall**
    - Allocate ROB/RS
    - Set RS tag to ROB#
    - Set Map Table entry to ROB# and clear “ready-in-ROB” bit
    - Read ready registers into RS (from either ROB or Regfile)
  - **X (execute)**
    - Free RS entry
      - No need to wait for W, because tag is from ROB instead of RS

# P6 Pipeline

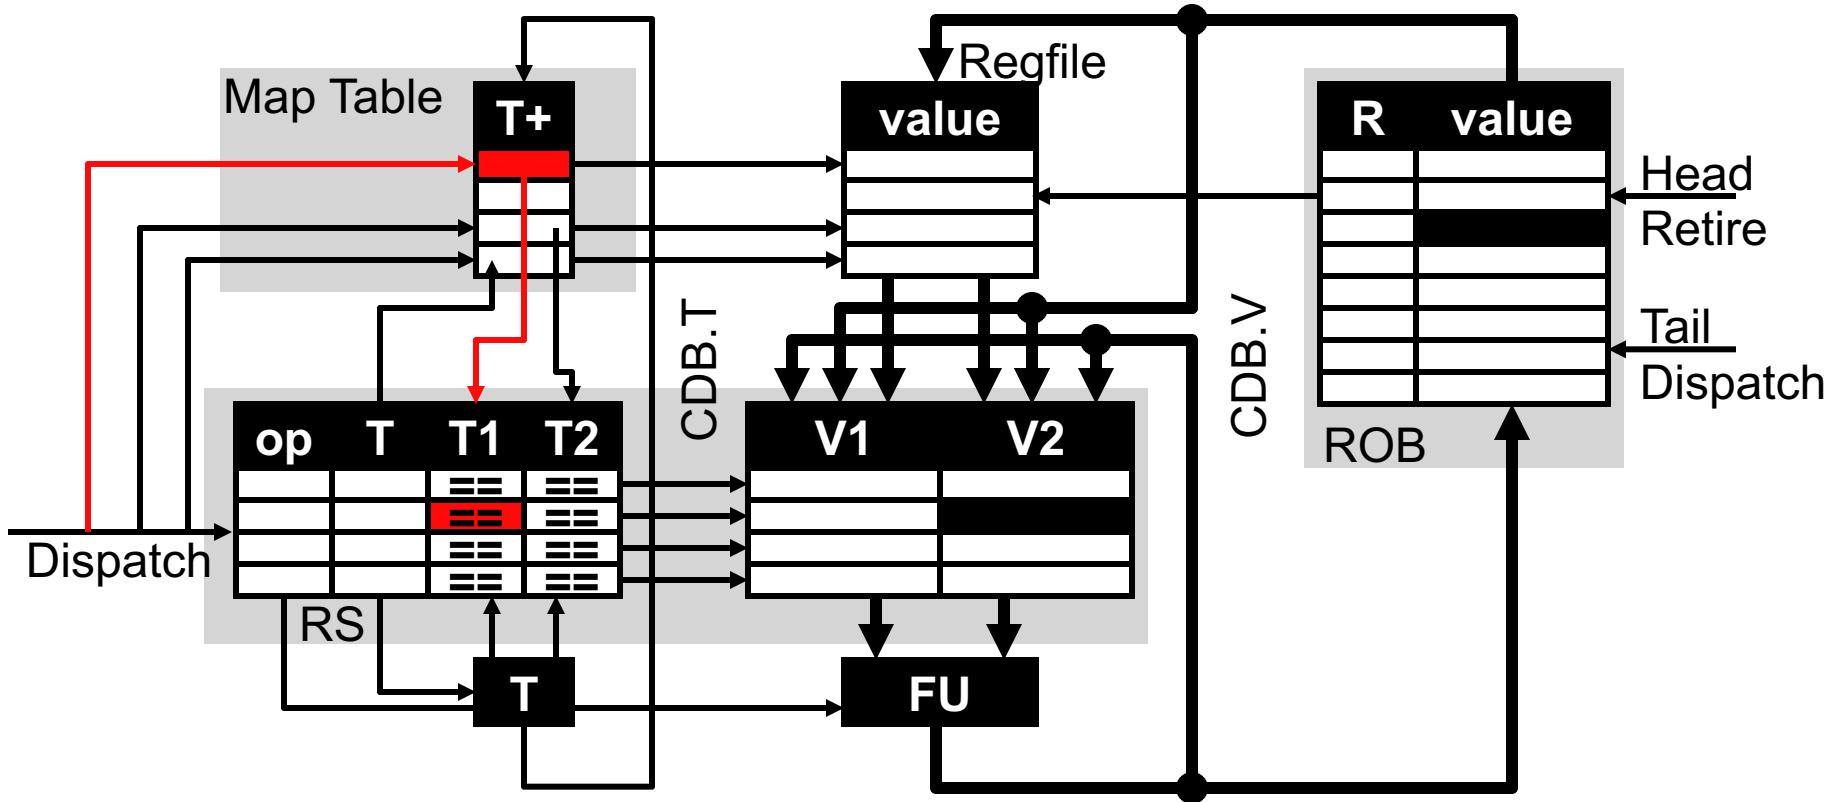
- **C (complete)**
  - Structural hazard (CDB)? **wait**
  - Write value into ROB entry for RS tag
  - If Map Table has same entry, set “ready-in-ROB” bit (+)
- **R (retire)**
  - Insn. at ROB head not complete ? **stall**
  - Handle any exceptions
    - Some go before instruction (branch mispredict, page fault) – why?
    - Some go after instruction (e.g., trap) – why?
  - ROB head value → Regfile
  - Free ROB entry

# P6 Dispatch (D) (1/2)



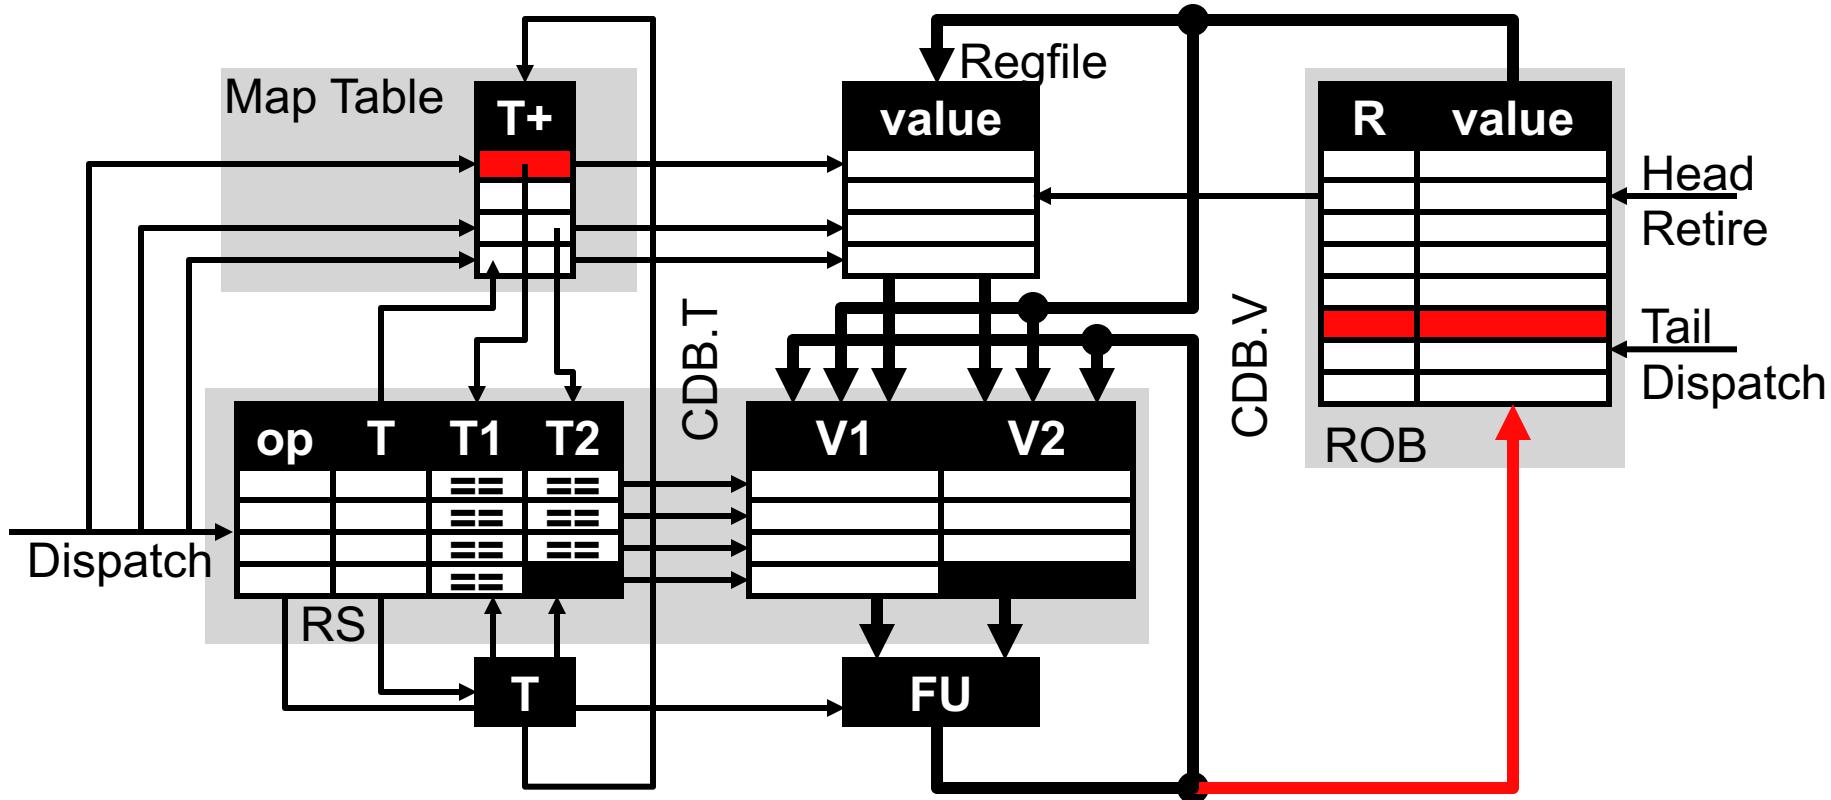
- RS/ROB full ? stall
- Allocate RS/ROB entries, assign ROB# to RS output tag
- Map Table entry set to ROB#, clear “ready-in-ROB”

# P6 Dispatch (D) (2/2)



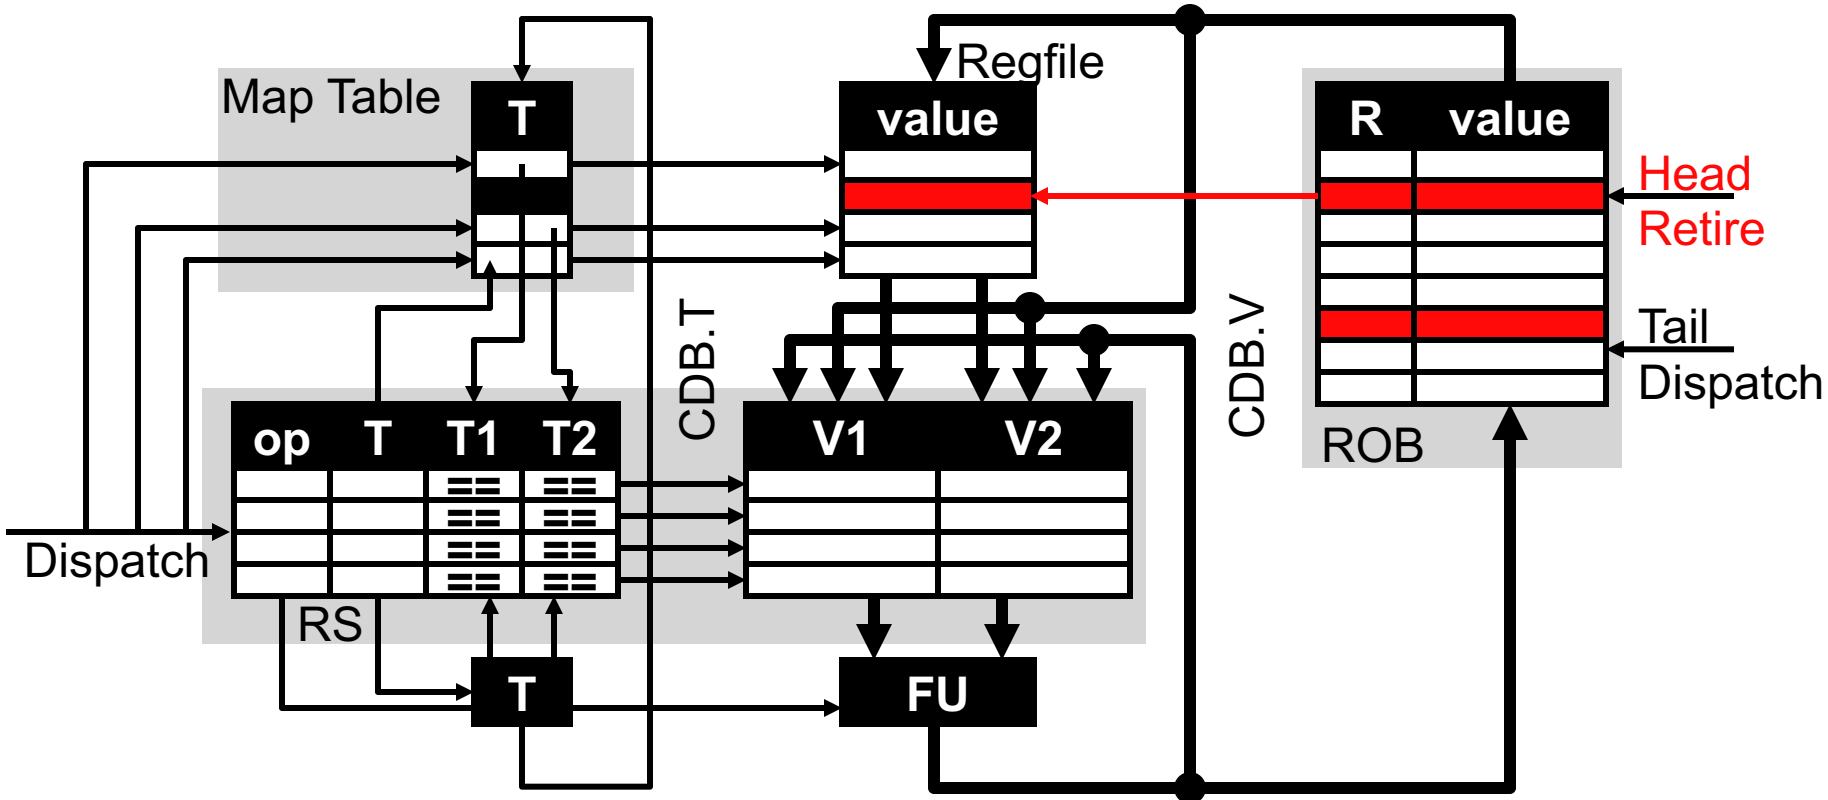
- Read tags for register inputs from Map Table
  - Tag==0 → value from Regfile (not shown)
  - Tag!=0 → Map Table tag to RS, Tag!=0+ → value from ROB

# P6 Complete (C)



- CDB busy ? stall : broadcast <value,tag> on CDB
- Result → ROB, if Map Table valid → “ready-in-ROB” bit
- If RS T1 or T2 matches, write CDB.V into RS slot

# P6 Retire (R)



- ROB head not complete ? stall : free ROB entry
- Write ROB head result to Regfile
- If still valid, clear Map Table entry

# P6: Cycle 1

ROB							
ht	#	Insn	R	V	S	X	C
ht	1	ldf X(r1), f1		f1			
	2	mulf f0, f1, f2					
	3	stf f2, Z(r1)					
	4	addi r1, 4, r1					
	5	ldf X(r1), f1					
	6	mulf f0, f1, f2					
	7	stf f2, Z(r1)					

Map Table	
Reg	T+
f0	
f1	ROB#1
f2	
r1	

CDB	
T	V

Reservation Stations								
#	FU	busy	op	T	T1	T2	V1	V2
1	ALU	no						
2	LD	yes	ldf	ROB#1			[r1]	
3	ST	no						
4	FP1	no						
5	FP2	no						

set ROB# tag

allocate

# P6: Cycle 2

ROB								
ht	#	Insn	R	V	S	X	C	
h	1	ldf X(r1), f1	f1		c2			
t	2	mult f0, f1, f2	f2					
	3	stf f2, Z(r1)						
	4	addi r1, 4, r1						
	5	ldf X(r1), f1						
	6	mult f0, f1, f2						
	7	stf f2, Z(r1)						

Map Table	
Reg	T+
f0	
f1	ROB#1
f2	ROB#2
r1	

CDB	
T	V

Reservation Stations								
#	FU	busy	op	T	T1	T2	V1	V2
1	ALU	no						
2	LD	yes	ldf	ROB#1				[r1]
3	ST	no						
4	FP1	yes	mult	ROB#2		ROB#1	[f0]	
5	FP2	no						

Map Table

Reg	T+
f0	
f1	ROB#1
f2	ROB#2
r1	

CDB

T	V

set ROB# tag

allocate

←

# P6: Cycle 3

ROB			R	V	S	X	C
ht	#	Insn					
h	1	ldf X(r1), f1	f1		c2	c3	
	2	mulf f0, f1, f2	f2				
t	3	stf f2, Z(r1)					
	4	addi r1, 4, r1					
	5	ldf X(r1), f1					
	6	mulf f0, f1, f2					
	7	stf f2, Z(r1)					

Map Table	
Reg	T+
f0	
f1	ROB#1
f2	ROB#2
r1	

CDB	
T	V

Reservation Stations								
#	FU	busy	op	T	T1	T2	V1	V2
1	ALU	no						
2	LD	no						
3	ST	yes	stf	ROB#3	ROB#2			[r1]
4	FP1	yes	mulf	ROB#2		ROB#1	[f0]	
5	FP2	no						

free  
allocate

# P6: Cycle 4

ROB							Map Table		CDB		
ht	#	Insn	R	V	S	X	C	Reg	T+	T	V
h	1	ldf X(r1), f1	f1	[f1]	c2	c3	c4	f0			
	2	mulf f0, f1, f2	f2		c4			f1	ROB#1+	ROB#1	[f1]
	3	stf f2, Z(r1)						f2	ROB#2		
t	4	addi r1, 4, r1	r1					r1	ROB#4		
	5	ldf X(r1), f1									
	6	mulf f0, f1, f2									
	7	stf f2, Z(r1)									

\ldf finished

- 1. set “ready-in-ROB” bit
- 2. write result to ROB
- 3. CDB broadcast

Reservation Stations									
#	FU	busy	op	T	T1	T2	V1	V2	
1	ALU	yes	add	ROB#4			[r1]		allocate
2	LD	no							
3	ST	yes	stf	ROB#3	ROB#2			[r1]	
4	FP1	yes	mulf	ROB#2		ROB#1	[f0]	CDB.V	ROB#1 ready grab CDB.V
5	FP2	no							

# P6: Cycle 5

ROB							Map Table		CDB		
ht	#	Insn	R	V	S	X	C	Reg	T+	T	V
	1	ldf X(r1), f1	f1	[f1]	c2	c3	c4	f0			
h	2	mulf f0, f1, f2	f2		c4	c5		f1	ROB#5		
	3	stf f2, Z(r1)						f2	ROB#2		
	4	addi r1, 4, r1	r1		c5			r1	ROB#4		
t	5	ldf X(r1), f1	f1								
	6	mulf f0, f1, f2									
	7	stf f2, Z(r1)									

ldf retires

1. write ROB result to regfile

Reservation Stations									
#	FU	busy	op	T	T1	T2	V1	V2	
1	ALU	yes	add	ROB#4			[r1]		
2	LD	yes	ldf	ROB#5		ROB#4			allocate
3	ST	yes	stf	ROB#3	ROB#2			[r1]	
4	FP1	no							free
5	FP2	no							

# P6: Cycle 6

ROB								
ht	#	Insn	R	V	S	X	C	
	1	ldf X(r1), f1	f1	[f1]	c2	c3	c4	
h	2	mulf f0, f1, f2	f2		c4	c5+		
	3	stf f2, Z(r1)						
	4	addi r1, 4, r1	r1		c5	c6		
	5	ldf X(r1), f1	f1					
t	6	mulf f0, f1, f2	f2					
	7	stf f2, Z(r1)						

Map Table	
Reg	T+
f0	
f1	ROB#5
f2	ROB#6
r1	ROB#4

CDB	
T	V

Reservation Stations								
#	FU	busy	op	T	T1	T2	V1	V2
1	ALU	no						
2	LD	yes	ldf	ROB#5		ROB#4		
3	ST	yes	stf	ROB#3	ROB#2			[r1]
4	FP1	yes	mulf	ROB#6		ROB#5	[f0]	
5	FP2	no						

free

allocate

# P6: Cycle 7

ROB								
ht	#	Insn	R	V	S	X	C	
	1	ldf X(r1), f1	f1	[f1]	c2	c3	c4	
h	2	mulf f0, f1, f2	f2		c4	c5+		
	3	stf f2, Z(r1)						
	4	addi r1, 4, r1	r1	[r1]	c5	c6	c7	
	5	ldf X(r1), f1	f1		c7			
t	6	mulf f0, f1, f2	f2					
	7	stf f2, Z(r1)						

stall D (no free STore RS)

Map Table	
Reg	T+
f0	
f1	ROB#5
f2	ROB#6
r1	ROB#4+

CDB	
T	V
ROB#4	[r1]

Reservation Stations								
#	FU	busy	op	T	T1	T2	V1	V2
1	ALU	no						
2	LD	yes	ldf	ROB#5		ROB#4		CDB.V
3	ST	yes	stf	ROB#3	ROB#2			[r1]
4	FP1	yes	mulf	ROB#6		ROB#5	[f0]	
5	FP2	no						

ROB#4 ready  
grab CDB.V

# P6: Cycle 8

ROB							Map Table		CDB		
ht	#	Insn	R	V	S	X	C	Reg	T+	T	V
	1	ldf X(r1), f1	f1	[f1]	c2	c3	c4	f0			
h	2	mulf f0, f1, f2	f2	[f2]	c4	c5+	c8	f1	ROB#5		
	3	stf f2, Z(r1)			c8			f2	ROB#6		
	4	addi r1, 4, r1	r1	[r1]	c5	c6	c7	r1	ROB#4+		
	5	ldf X(r1), f1	f1		c7	c8					
t	6	mulf f0, f1, f2	f2								
	7	stf f2, Z(r1)									

stall R for addi (in-order)  
 ROB#2 not in MapTable f2,  
 don't set "ready-in-ROB"

Reservation Stations									
#	FU	busy	op	T	T1	T2	V1	V2	
1	ALU	no							
2	LD	no							
3	ST	yes	stf	ROB#3	ROB#2		[f2]	[r1]	
4	FP1	yes	mulf	ROB#6		ROB#5	[f0]		
5	FP2	no							

ROB#2 ready  
 grab CDB.V

# P6: Cycle 9

ROB								
ht	#	Insn	R	V	S	X	C	
	1	ldf X(r1), f1	f1	[f1]	c2	c3	c4	
	2	mulf f0, f1, f2	f2	[f2]	c4	c5+	c8	
h	3	stf f2, Z(r1)			c8	c9		
	4	addi r1, 4, r1	r1	[r1]	c5	c6	c7	
	5	ldf X(r1), f1	f1	[f1]	c7	c8	c9	
	6	mulf f0, f1, f2	f2		c9			
t	7	stf f2, Z(r1)						

Map Table	
Reg	T+
f0	
f1	ROB#5+
f2	ROB#6
r1	ROB#4+

CDB	
T	V
ROB#5	[f1]

retire mulf

all pipe stages active at once!

Reservation Stations								
#	FU	busy	op	T	T1	T2	V1	V2
1	ALU	no						
2	LD	no						
3	ST	yes	stf	ROB#7	ROB#6			ROB#4 . V
4	FP1	yes	mulf	ROB#6		ROB#5	[f0]	CDB . V
5	FP2	no						

free, re-allocate  
ROB#5 ready  
grab CDB.V

# P6: Cycle 10

ROB									Map Table		CDB	
ht	#	Insn	R	V	S	X	C	Reg	T+	T	V	
	1	ldf X(r1), f1	f1	[f1]	c2	c3	c4	f0				
	2	mulf f0, f1, f2	f2	[f2]	c4	c5+	c8	f1	ROB#5+			
h	3	stf f2, Z(r1)			c8	c9	c10	f2	ROB#6			
	4	addi r1, 4, r1	r1	[r1]	c5	c6	c7	r1	ROB#4+			
	5	ldf X(r1), f1	f1	[f1]	c7	c8	c9					
	6	mulf f0, f1, f2	f2		c9	c10						
t	7	stf f2, Z(r1)										

Reservation Stations									
#	FU	busy	op	T	T1	T2	V1	V2	
1	ALU	no							
2	LD	no							
3	ST	yes	stf	ROB#7	ROB#6			ROB#4 . V	
4	FP1	no							free
5	FP2	no							

# P6: Cycle 11

ROB

ht	#	Insn	R	V	S	X	C
	1	ldf X(r1), f1	f1	[f1]	c2	c3	c4
	2	mulf f0, f1, f2	f2	[f2]	c4	c5	c8
	3	stf f2, Z(r1)			c8	c9	c10
h	4	addi r1, 4, r1	r1	[r1]	c5	c6	c7
	5	ldf X(r1), f1	f1	[f1]	c7	c8	c9
	6	mulf f0, f1, f2	f2		c9	c10	
t	7	stf f2, Z(r1)					

Map Table

Reg	T+
f0	
f1	ROB#5+
f2	ROB#6
r1	ROB#4+

CDB

T	V

retire stf

Reservation Stations

#	FU	busy	op	T	T1	T2	V1	V2
1	ALU	no						
2	LD	no						
3	ST	yes	stf	ROB#7	ROB#6			ROB#4 . V
4	FP1	no						
5	FP2	no						

# Precise State in P6

- Point of ROB is maintaining **precise state**
  - How does that work?
    1. Wait until last good insn. retires, first bad insn. at ROB head
    2. Zero (0) contents of ROB, RS, and Map Table
    3. Start over
  - Works because zero (0) means the right thing...
    - 0 in ROB/RS → entry is empty
    - Tag == 0 in Map Table → register is in Regfile
  - ...and because Regfile and L1-D writes take place at R
  - Example: page fault in first **stf**

# P6: Cycle 9 (with precise state)

ROB						
ht	#	Insn	R	V	S	X
	1	ldf X(r1), f1	f1	[f1]	c2	c3
	2	mulf f0, f1, f2	f2	[f2]	c4	c5+
h	3	stf f2, Z(r1)			c8	c9
	4	addi r1, 4, r1	r1	[r1]	c5	c6
	5	ldf X(r1), f1	f1	[f1]	c7	c8
	6	mulf f0, f1, f2	f2		c9	
t	7	stf f2, Z(r1)				

Map Table	
Reg	T+
f0	
f1	ROB#5+
f2	ROB#6
r1	ROB#4+

CDB	
T	V
ROB#5	[f1]

PAGE FAULT

Reservation Stations								
#	FU	busy	op	T	T1	T2	V1	V2
1	ALU	no						
2	LD	no						
3	ST	yes	stf	ROB#7	ROB#6			ROB#4 . V
4	FP1	yes	mulf	ROB#6		ROB#5	[f0]	CDB . V
5	FP2	no						

# P6: Cycle 10 (with precise state)

ROB							Map Table		CDB		
ht	#	Insn	R	V	S	X	C	Reg	T+	T	V
	1	<code>ldf X(r1), f1</code>	f1	[f1]	c2	c3	c4	f0			
	2	<code>mulf f0, f1, f2</code>	f2	[f2]	c4	c5+	c8	f1			
	3	<code>stf f2, Z(r1)</code>						f2			
	4	<code>addi r1, 4, r1</code>						r1			
	5	<code>ldf X(r1), f1</code>									
	6	<code>mulf f0, f1, f2</code>									
	7	<code>stf f2, Z(r1)</code>									

faulting insn at ROB head?  
CLEAR EVERYTHING  
set fetch PC to fault handler

Reservation Stations									
#	FU	busy	op	T	T1	T2	V1	V2	
1	ALU	no							
2	LD	no							
3	ST	no							
4	FP1	no							
5	FP2	no							

# P6: Cycle 11 (with precise state)

ROB							
ht	#	Insn	R	V	S	X	C
	1	ldf X(r1), f1	f1	[f1]	c2	c3	c4
	2	mulf f0, f1, f2	f2	[f2]	c4	c5+	c8
ht	3	stf f2, Z(r1)					
	4	addi r1, 4, r1					
	5	ldf X(r1), f1					
	6	mulf f0, f1, f2					
	7	stf f2, Z(r1)					

Map Table	
Reg	T+
f0	
f1	
f2	
r1	

CDB	
T	V

PF handler done?  
 CLEAR EVERYTHING  
 iret fetch PC to faulting insn.

Reservation Stations									
#	FU	busy	op	T	T1	T2	V1	V2	
1	ALU	no							
2	LD	no							
3	ST	yes	stf	ROB#3			[f4]	[r1]	
4	FP1	no							
5	FP2	no							

# P6: Cycle 12 (with precise state)

ROB							
ht	#	Insn	R	V	S	X	C
	1	ldf X(r1), f1	f1	[f1]	c2	c3	c4
	2	mulf f0, f1, f2	f2	[f2]	c4	c5+	c8
h	3	stf f2, Z(r1)			c12		
t	4	addi r1, 4, r1	r1				
	5	ldf X(r1), f1					
	6	mulf f0, f1, f2					
	7	stf f2, Z(r1)					

Map Table	
Reg	T+
f0	
f1	
f2	
r1	ROB#4

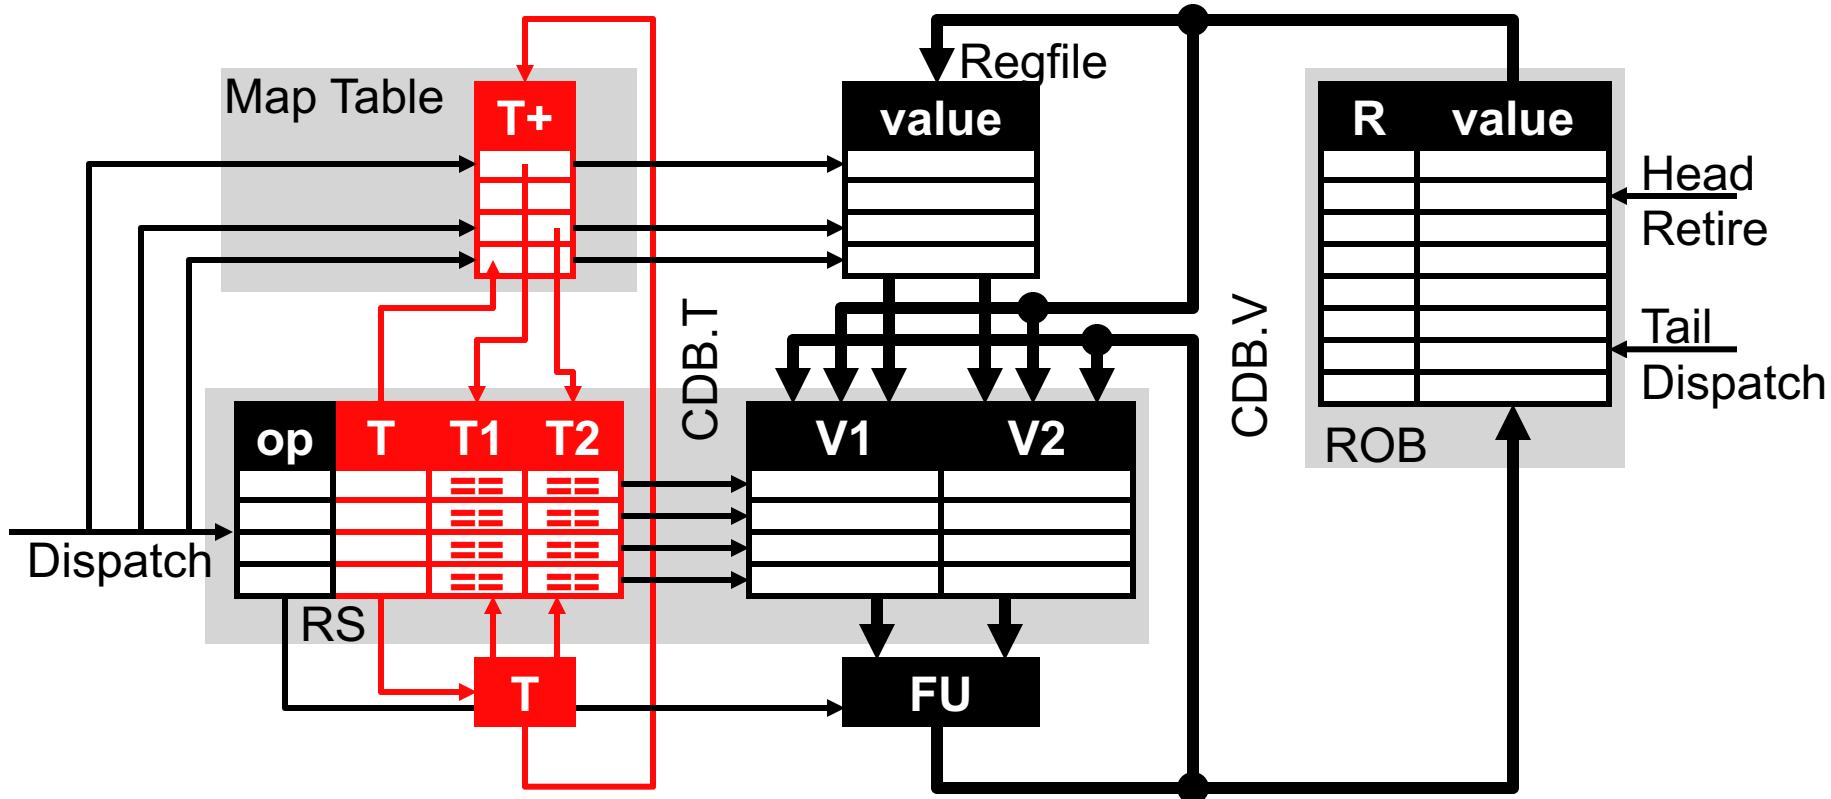
CDB	
T	V

Reservation Stations								
#	FU	busy	op	T	T1	T2	V1	V2
1	ALU	yes	addi	ROB#4			[r1]	
2	LD	no						
3	ST	yes	stf	ROB#3			[f4]	[r1]
4	FP1	no						
5	FP2	no						

# P6 Performance

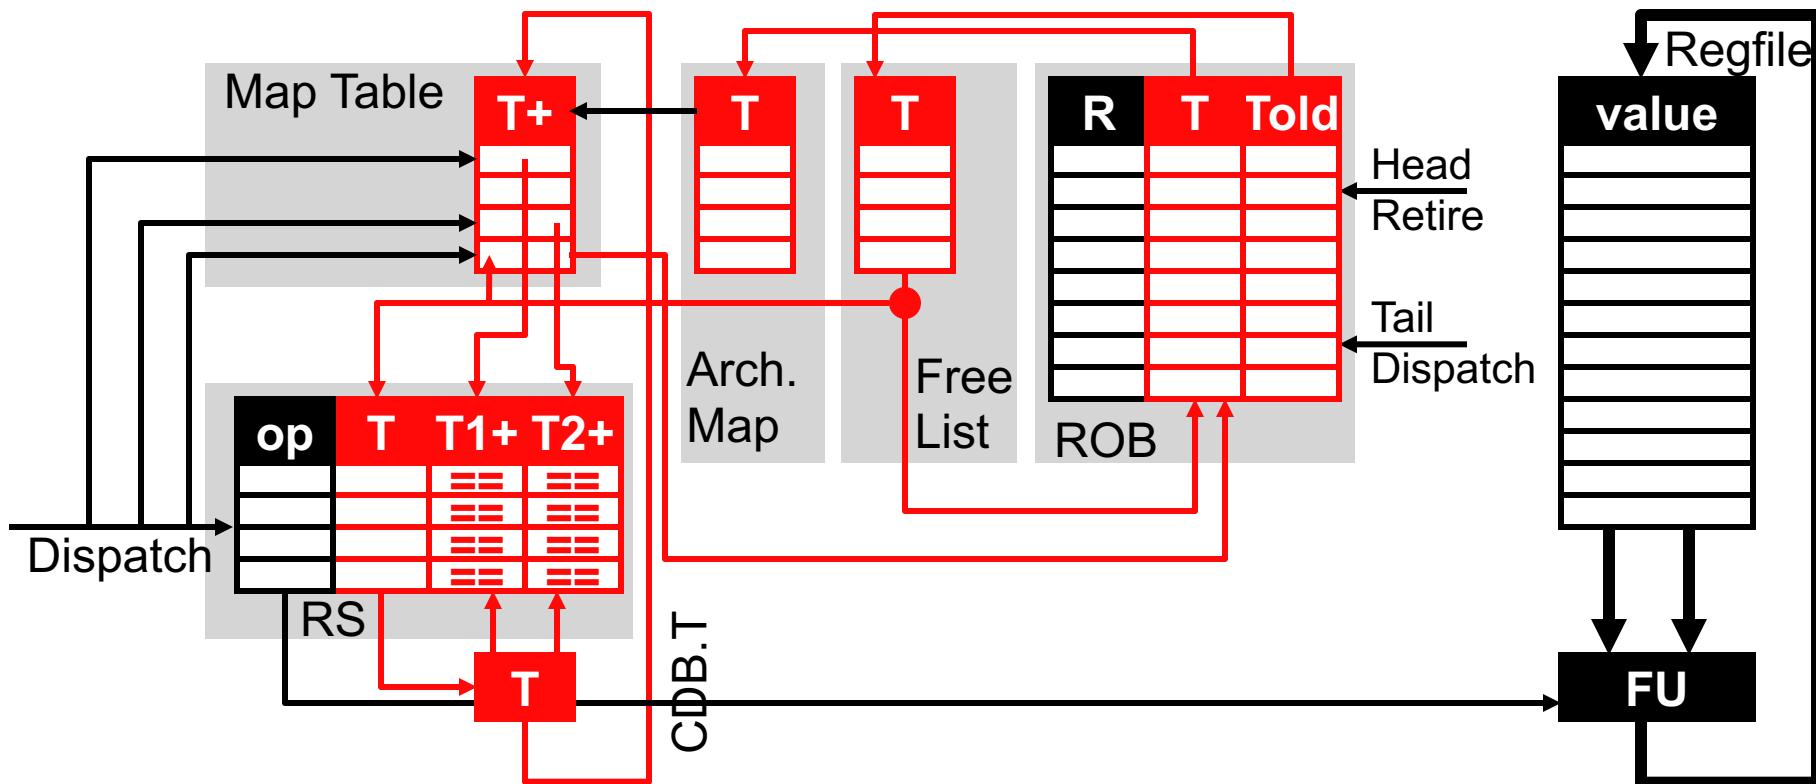
- In other words: what is the cost of precise state?
  - + In general: same performance as “plain” Tomasulo
    - ROB is not a performance device
    - Maybe a little better (RS freed earlier → fewer struct hazards)
  - Unless ROB is too small
    - In which case ROB struct hazards become a problem
  - Rules of thumb for ROB size
    - At least  $N$  (width) \* number of pipe stages between D and R
    - At least  $N * t_{hit-L2}$
    - Can add a factor of 2 to both if you want
    - What is the rationale behind these?

# The Problem with P6



- Problem for high performance implementations
  - Too much **value movement** (Regfile/ROB→RS→ROB→Regfile)
  - Multi-input muxes, long buses, slow clock

# MIPS R10K: Alternative Implementation



- One big **physical register file** holds all data - no copies
    - + Register file close to FUs → small and fast data path
    - ROB and RS “on the side” used only for control and tags

# Register Renaming in R10K

- Architectural register file? Gone
- Physical register file holds all values
  - #physical registers = #architectural registers + #ROB entries
  - Map architectural registers to physical registers
  - No WAW or WAR hazards (physical regs. replace RS values)
- Fundamental change to map table
  - Mappings cannot be 0 (no architectural register file)
- Explicit free list tracks unallocated physical regs.
  - ROB returns physical regs. to free list

# Register Renaming

- Example
  - Names: r1,r2,r3 Locations: p1,p2,p3,p4,p5,p6,p7
  - Original:  $r1 \rightarrow p1$ ,  $r2 \rightarrow p2$ ,  $r3 \rightarrow p3$ ,  $p4-p7$  are “free”

MapTable

r1	r2	r3
p1	p2	p3
p4	p2	p3
p4	p2	p5
p4	p2	p6

FreeList

p4 , p5 , p6 , p7
p5 , p6 , p7
p6 , p7
p7

Original insns.

add r2 , r3 , r1
sub r2 , r1 , r3
mul r2 , r3 , r3
div r1 , 4 , r1

Renamed insns.

add p2 , p3 , p4
sub p2 , p4 , p5
mul p2 , p5 , p6
div p4 , 4 , p7

# Register Renaming

- Example
  - Names: r1,r2,r3 Locations: p1,p2,p3,p4,p5,p6,p7
  - Original:  $r1 \rightarrow p1$ ,  $r2 \rightarrow p2$ ,  $r3 \rightarrow p3$ ,  $p4-p7$  are “free”

MapTable	r1	r2	r3
p1	p2	p3	
p4	p2	p3	
p4	p2	p5	
p4	p2	p6	
p7	p2	p6	

FreeList
p4 , p5 , p6 , p7
p5 , p6 , p7
p6 , p7
p7

Original insns.

add r2,r3,r1  
sub r2,r1,r3  
mul r2,r3,r3  
div r1,4,r1  
add r1,r3,r2

Renamed insns.

add p2,p3,p4  
sub p2,p4,p5  
mul p2,p5,p6  
div p4,4,p7  
add p7,p6,???

• Question: how is the last *add* renamed?

- We are out of free physical registers
- Real question: how/when are physical registers freed?

# Physical Register Reclamation

- P6
  - No need to free speculative (“in-flight”) values explicitly
  - Temporary storage comes with ROB entry
- R10K
  - Can’t free physical regs. when insn. retires
    - Younger insns. likely depend on it
  - But...
    - Can free physical reg. previously mapped to same logical reg.
    - Why?

# Freeing Registers in R10K

MapTable

r1	r2	r3
p1	p2	p3
p4	p2	p3
p4	p2	p5
p4	p2	p6
p7	p2	p6

FreeList

p4, p5, p6, p7
p5, p6, p7
p6, p7
p7
p1

Original insns.

```
add r2, r3, r1
sub r2, r1, r3
mul r2, r3, r3
div r1, 4, r1
add r1, r3, r2
```

Renamed insns.

```
add p2, p3, p4
sub p2, p4, p5
mul p2, p5, p6
div p4, 4, p7
add p7, p6, p1
```

- When ***add*** retires, free p1
- When ***sub*** retires, free p3
- When ***mul*** retires, free p5
- When ***div*** retires, free p4

Always OK to free ***old*** mapping

# R10K Data Structures

- New tags (again)
  - P6: ROB# → R10K: PR#
- ROB
  - **T**: physical register corresponding to insn's logical output
  - **Told**: physical register previously mapped to insn's logical output
- RS
  - **T, T1, T2**: output, input physical registers
- Map Table
  - **T+**: PR# (never empty) + “ready” bit
- Architectural Map Table
  - **T**: PR# (never empty)
- Free List
  - **T**: PR#

No values in ROB, RS, or on CDB

# R10K Data Structures

ROB						
ht	#	Insn	T	Told	S	X
	1	ldf X(r1), f1				
	2	mulf f0, f1, f2				
	3	stf f2, Z(r1)				
	4	addi r1, 4, r1				
	5	ldf X(r1), f1				
	6	mulf f0, f1, f2				
	7	stf f2, Z(r1)				

Map Table	
Reg	T+
f0	PR#1+
f1	PR#2+
f2	PR#3+
r1	PR#4+

Arch. Map	
Reg	T+
f0	PR#1
f1	PR#2
f2	PR#3
r1	PR#4

Free List	
PR#5, PR#6, PR#7, PR#8	

CDB	
T	

Reservation Stations						
#	FU	busy	op	T	T1	T2
1	ALU	no				
2	LD	no				
3	ST	no				
4	FP1	no				
5	FP2	no				

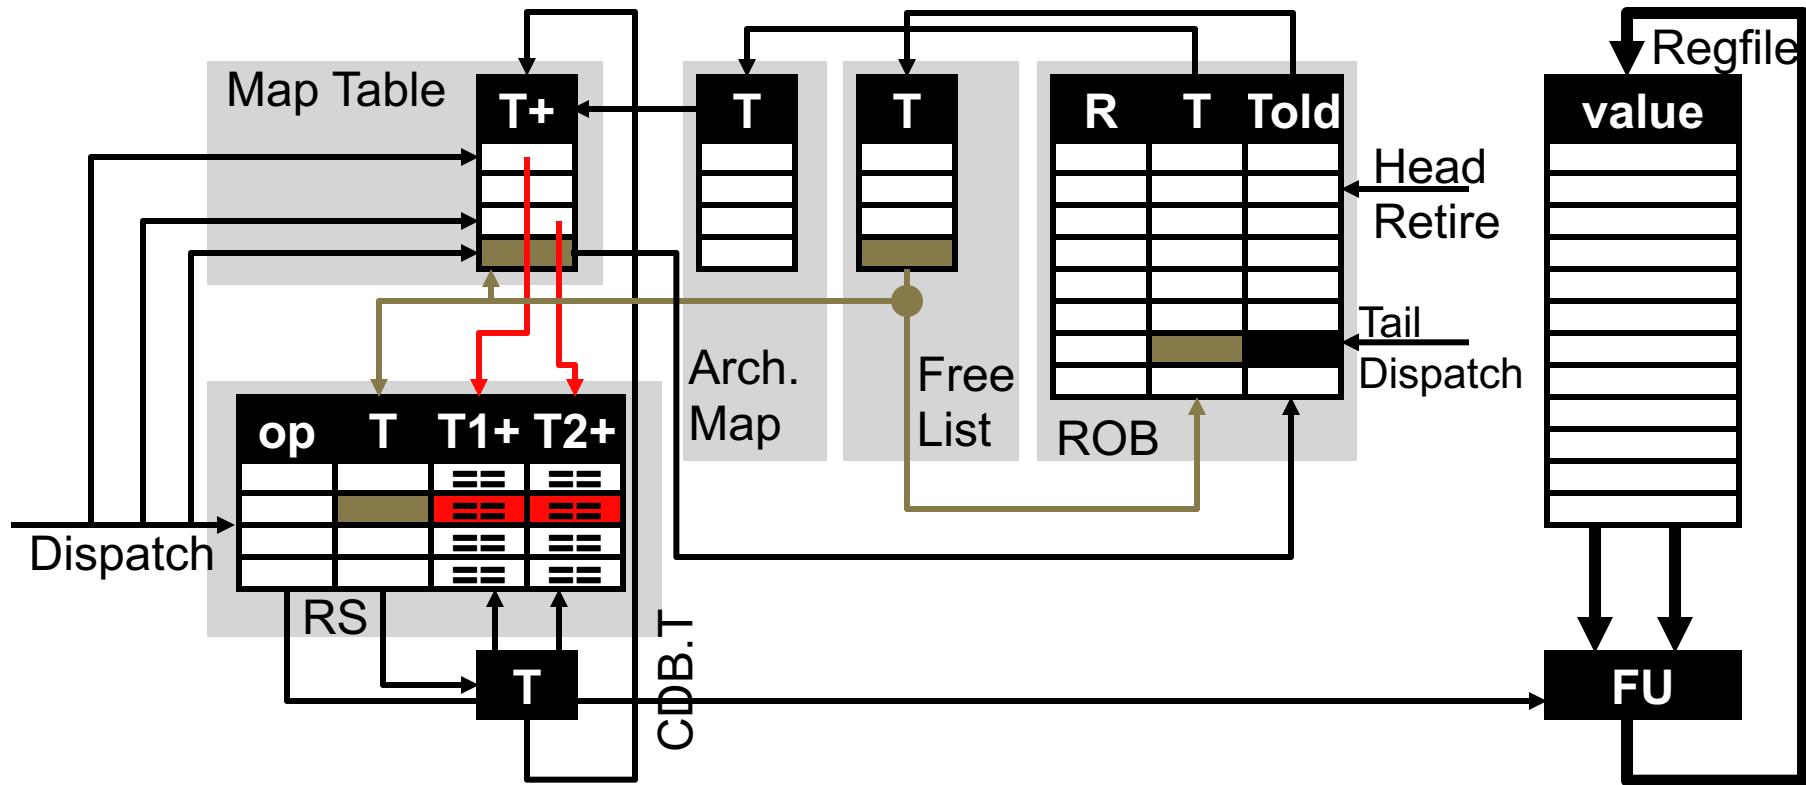
Notice I: no values anywhere

Notice II: MapTable is never empty

# R10K Pipeline

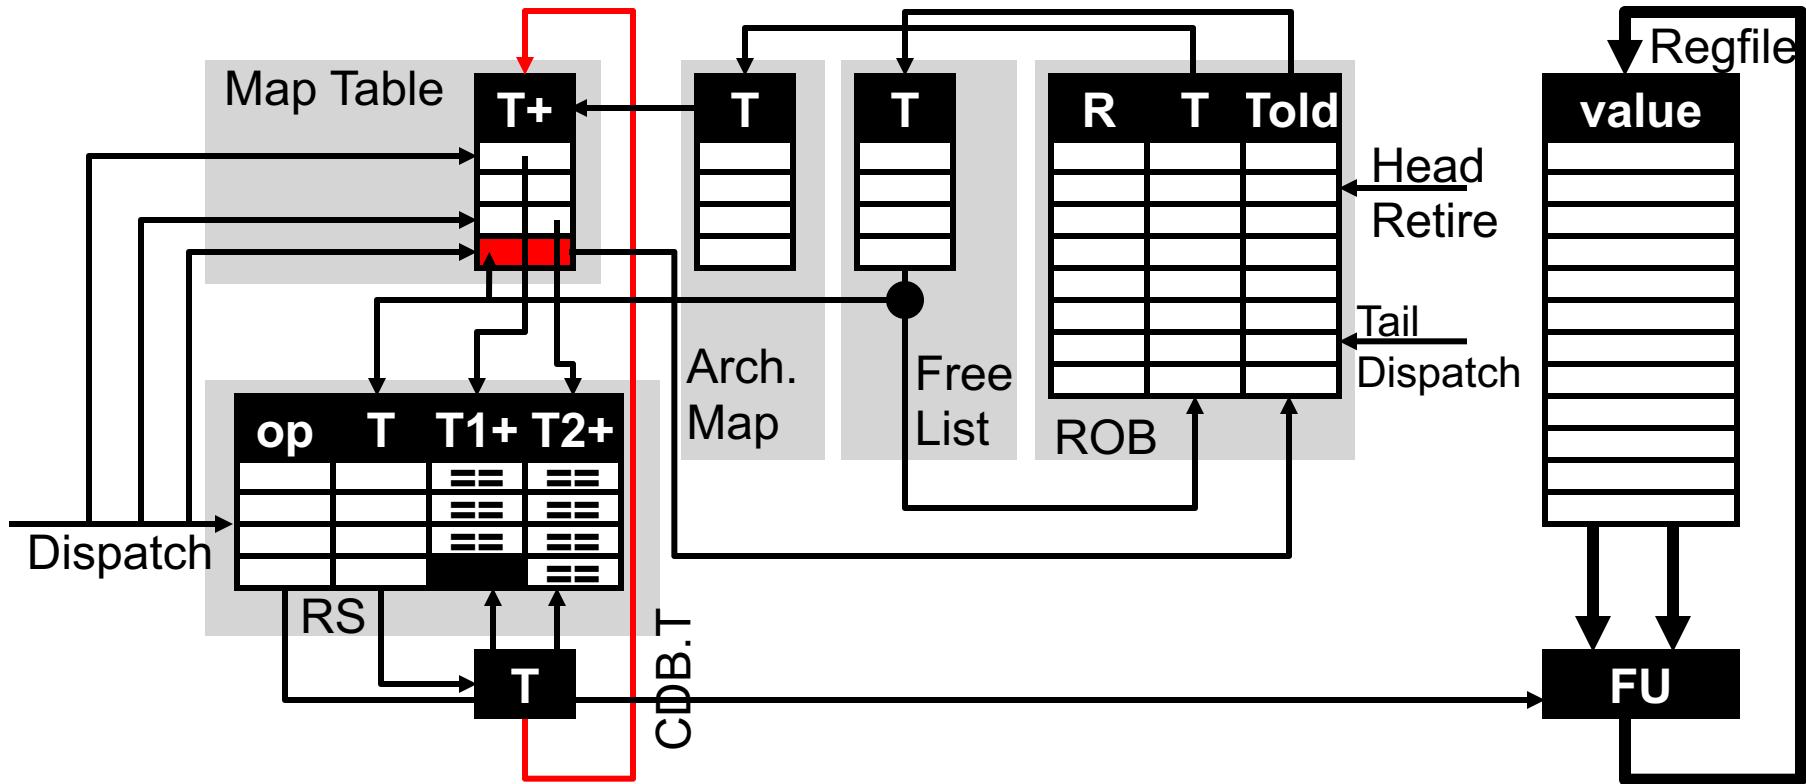
- R10K pipeline structure: F, **D**, S, X, C, **R**
  - **D (dispatch)**
    - Structural hazard (RS, ROB, **physical registers**) ? stall
    - Allocate RS, ROB, and new physical register (T)
    - **Record previously mapped physical register (Told)**
  - **C (complete)**
    - Write destination physical register
  - **R (retire)**
    - ROB head not complete ? stall
    - Handle any exceptions
    - Free ROB entry
    - **Free previous physical register (Told)**
    - **Record committed physical register (T)**

# R10K Dispatch (D)



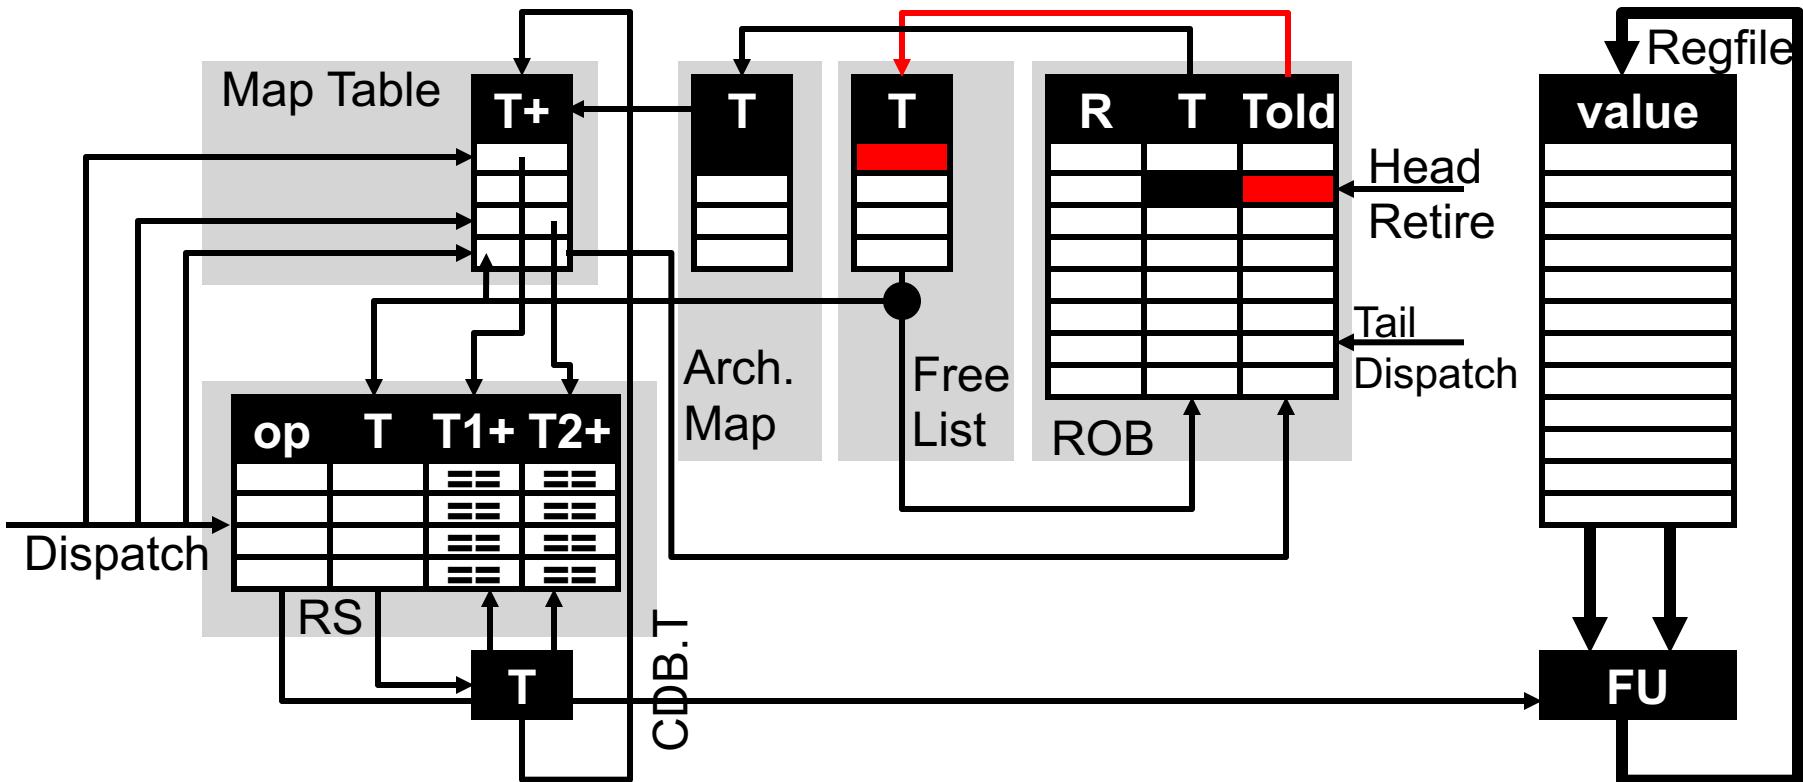
- Read preg (physical register) tags for input registers, store in RS
- Read preg tag for output register, store in ROB (Told)
- Allocate new preg (free list) for output reg, store in RS, ROB, Map Table

# R10K Complete (C)



- Set insn's output register ready bit in map table
- Set ready bits for matching input tags in RS

# R10K Retire (R)



- Return Told of ROB head to free list
- Record T of ROB head in architectural map table

# R10K: Cycle 1

ROB						
ht	#	Insn	T	Told	S	X
ht	1	ldf X(r1), f1	PR#5	PR#2		
	2	mulf f0, f1, f2				
	3	stf f2, Z(r1)				
	4	addi r1, 4, r1				
	5	ldf X(r1), f1				
	6	mulf f0, f1, f2				
	7	stf f2, Z(r1)				

Map Table	
Reg	T+
f0	PR#1+
f1	PR#5
f2	PR#3+
r1	PR#4+

Free List	
PR#5	, PR#6,
PR#7	, PR#8

Arch. Map	
Reg	T+
f0	PR#1
f1	PR#2
f2	PR#3
r1	PR#4

CDB
T

Reservation Stations						
#	FU	busy	op	T	T1	T2
1	ALU	no				
2	LD	yes	ldf	PR#5		PR#4+
3	ST	no				
4	FP1	no				
5	FP2	no				

Allocate new preg (PR#5) to f1

Remember old preg mapped to f1 (PR#2) in ROB

# R10K: Cycle 2

ROB							
ht	#	Insn	T	Told	S	X	C
h	1	ldf X(r1), f1	PR#5	PR#2	c2		
t	2	mulf f0, f1, f2	PR#6	PR#3			
	3	stf f2, Z(r1)					
	4	addi r1, 4, r1					
	5	ldf X(r1), f1					
	6	mulf f0, f1, f2					
	7	stf f2, Z(r1)					

Map Table	
Reg	T+
f0	PR#1+
f1	PR#5
f2	PR#6
r1	PR#4+

Free List	
PR#6	, PR#7 , PR#8

Arch. Map	
Reg	T+
f0	PR#1
f1	PR#2
f2	PR#3
r1	PR#4

CDB
T

Reservation Stations						
#	FU	busy	op	T	T1	T2
1	ALU	no				
2	LD	yes	ldf	PR#5		PR#4+
3	ST	no				
4	FP1	yes	mulf	PR#6	PR#1+ PR#5	
5	FP2	no				

Allocate new preg (PR#6) to f2

Remember old preg mapped to f3 (PR#3) in ROB

# R10K: Cycle 3

ROB							Map Table	Arch. Map		
ht	#	Insn	T	Told	S	X	Reg	T+	Reg	T+
h	1	ldf X(r1), f1	PR#5	PR#2	c2	c3	f0	PR#1+	f0	PR#1
	2	mulf f0, f1, f2	PR#6	PR#3			f1	PR#5	f1	PR#2
t	3	stf f2, Z(r1)					f2	PR#6	f2	PR#3
	4	addi r1, 4, r1					r1	PR#4+	r1	PR#4
	5	ldf X(r1), f1								
	6	mulf f0, f1, f2								
	7	stf f2, Z(r1)								

Reservation Stations							Free List	CDB
#	FU	busy	op	T	T1	T2		
1	ALU	no						
2	LD	no						
3	ST	yes	stf		PR#6	PR#4+		
4	FP1	yes	mulf	PR#6	PR#1+	PR#5		
5	FP2	no						

Stores are not allocated pregs

Free

# R10K: Cycle 4

ROB							
ht	#	Insn	T	Told	S	X	C
h	1	ldf X(r1), f1	PR#5	PR#2	c2	c3	c4
	2	mulf f0, f1, f2	PR#6	PR#3	c4		
	3	stf f2, Z(r1)					
t	4	addi r1, 4, r1	PR#7	PR#4			
	5	ldf X(r1), f1					
	6	mulf f0, f1, f2					
	7	stf f2, Z(r1)					

Map Table	
Reg	T+
f0	PR#1+
f1	PR#5+
f2	PR#6
r1	PR#7

Arch. Map	
Reg	T+
f0	PR#1
f1	PR#2
f2	PR#3
r1	PR#4

Free List	
PR#7	, PR#8

Reservation Stations						
#	FU	busy	op	T	T1	T2
1	ALU	yes	addi	PR#7	PR#4+	
2	LD	no				
3	ST	yes	stf		PR#6	PR#4+
4	FP1	yes	mulf	PR#6	PR#1+	PR#5+
5	FP2	no				

ldf completes  
set MapTable ready bit

Match PR#5 tag from CDB & issue

# R10K: Cycle 5

ROB						
ht	#	Insn	T	Told	S	X
	1	ldf X(r1), f1	PR#5 PR#2	c2 c3 c4		
h	2	mulf f0, f1, f2	PR#6 PR#3	c4 c5		
	3	stf f2, Z(r1)				
	4	addi r1, 4, r1	PR#7 PR#4	c5		
t	5	ldf X(r1), f1	PR#8 PR#5			
	6	mulf f0, f1, f2				
	7	stf f2, Z(r1)				

Map Table	
Reg	T+
f0	PR#1+
f1	PR#8
f2	PR#6
r1	PR#7

Arch. Map	
Reg	T+
f0	PR#1
f1	PR#5
f2	PR#3
r1	PR#4

Free List
PR#8, PR#2

CDB
T

Reservation Stations						
#	FU	busy	op	T	T1	T2
1	ALU	yes	addi	PR#7	PR#4+	
2	LD	yes	ldf	PR#8		PR#7
3	ST	yes	stf		PR#6	PR#4+
4	FP1	no				
5	FP2	no				

Free

ldf retires  
Return PR#2 to free list  
Record PR#5 in Arch map

# Precise State in R10K

- Precise state is more difficult in R10K
  - Physical registers are written out-of-order (at C)
  - Roll back the Map Table, Arch Table, Free List
    - “free” written registers and “restore” old ones
- Two ways of restoring Map Table and Free List
  - Option I: serial rollback using  $T$ ,  $T_{old}$  ROB fields
    - ± Slow, but simple
  - Option II: single-cycle restoration from some checkpoint
    - ± Fast, but checkpoints are expensive
  - Modern processor compromise: **make common case fast**
    - Checkpoint only (low-confidence) branches (frequent rollbacks)
    - Serial recovery for page-faults and interrupts (rare rollbacks)

# R10K: Cycle 5 (with precise state)

ROB							
ht	#	Insn	T	Told	S	X	C
	1	ldf X(r1), f1	PR#5	PR#2	c2	c3	c4
h	2	mulf f0, f1, f2	PR#6	PR#3	c4	c5	
	3	stf f2, Z(r1)					
	4	addi r1, 4, r1	PR#7	PR#4	c5		
t	5	ldf X(r1), f1	PR#8	PR#5			
	6	mulf f0, f1, f2					
	7	stf f2, Z(r1)					

Map Table	
Reg	T+
f0	PR#1+
f1	PR#8
f2	PR#6
r1	PR#7

Free List	
PR#8	PR#2

CDB	
T	

Reservation Stations						
#	FU	busy	op	T	T1	T2
1	ALU	yes	addi	PR#7	PR#4+	
2	LD	yes	ldf	PR#8		PR#7
3	ST	yes	stf		PR#6	PR#4+
4	FP1	no				
5	FP2	no				

**undo insns 3-5**  
 (doesn't matter why)  
 use serial rollback

# R10K: Cycle 6 (with precise state)

ROB							
ht	#	Insn	T	Told	S	X	C
	1	ldf X(r1), f1	PR#5	PR#2	c2	c3	c4
h	2	mulf f0, f1, f2	PR#6	PR#3	c4	c5	
	3	stf f2, Z(r1)					
t	4	addi r1, 4, r1	PR#7	PR#4	c5		
	5	ldf X(r1), f1	PR#8	PR#5			
	6	mulf f0, f1, f2					
	7	stf f2, Z(r1)					

Map Table	
Reg	T+
f0	PR#1+
f1	PR#5+
f2	PR#6
r1	PR#7

CDB	
T	

Free List	
PR#2	PR#8

Reservation Stations						
#	FU	busy	op	T	T1	T2
1	ALU	yes	addi	PR#7	PR#4+	
2	LD	no				
3	ST	yes	stf		PR#6	PR#4+
4	FP1	no				
5	FP2	no				

undo ldf (ROB#5)

1. free RS
2. free T (PR#8), return to FreeList
3. restore MT[f1] to Told (PR#5)
4. free ROB#5

insns may execute during rollback  
(not shown)

# R10K: Cycle 7 (with precise state)

ROB							
ht	#	Insn	T	Told	S	X	C
	1	ldf X(r1), f1	PR#5	PR#2	c2	c3	c4
h	2	mulf f0, f1, f2	PR#6	PR#3	c4	c5	
t	3	stf f2, Z(r1)					
	4	addi r1, 4, r1	PR#7	PR#4	c5		
	5	ldf X(r1), f1					
	6	mulf f0, f1, f2					
	7	stf f2, Z(r1)					

Map Table	
Reg	T+
f0	PR#1+
f1	PR#5+
f2	PR#6
r1	PR#4+

CDB	
T	

Free List	
PR#2, PR#8, PR#7	

Reservation Stations						
#	FU	busy	op	T	T1	T2
1	ALU	no				
2	LD	no				
3	ST	yes	stf		PR#6	PR#4+
4	FP1	no				
5	FP2	no				

undo addi (ROB#4)

1. free RS
2. free T (PR#7), return to FreeList
3. restore MT[r1] to Told (PR#4)
4. free ROB#4

# R10K: Cycle 8 (with precise state)

ROB							
ht	#	Insn	T	Told	S	X	C
	1	ldf X(r1), f1	PR#5	PR#2	c2	c3	c4
ht	2	mulf f0, f1, f2	PR#6	PR#3	c4	c5	
	3	stf f2, Z(r1)					
	4	addi r1, 4, r1					
	5	ldf X(r1), f1					
	6	mulf f0, f1, f2					
	7	stf f2, Z(r1)					

Map Table	
Reg	T+
f0	PR#1+
f1	PR#5+
f2	PR#6
r1	PR#4+

Free List	
PR#2, PR#8,	PR#7

CDB	
T	

Reservation Stations						
#	FU	busy	op	T	T1	T2
1	ALU	no				
2	LD	no				
3	ST	no				
4	FP1	no				
5	FP2	no				

undo stf (ROB#3)

1. free RS
2. free ROB#3
3. no registers to restore/free
4. how is L1-D write undone?

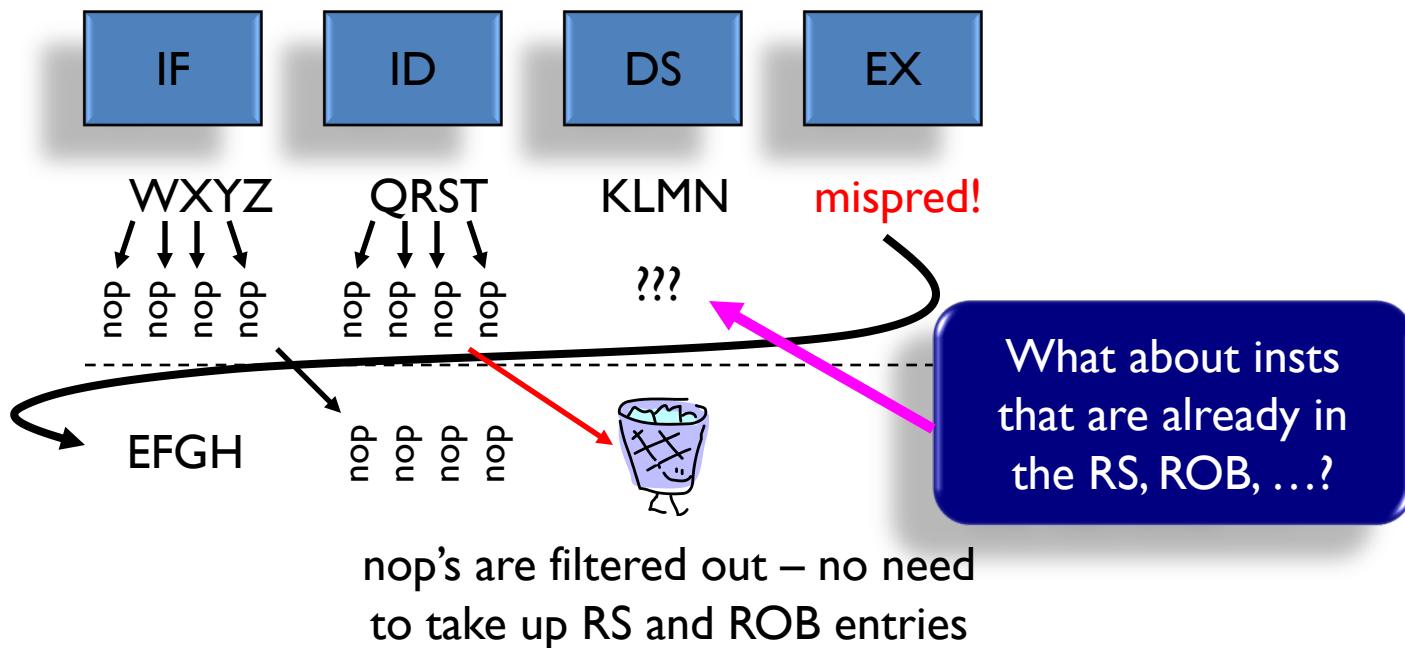
# P6 vs. R10K (Renaming)

Feature	P6	R10K
Value storage	ARF,ROB,RS	PRF
Register read	@D: ARF/ROB → RS	@S: PRF → FU
Register write	@R: ROB → ARF	@C: FU → PRF
Speculative value free	@R: automatic (ROB)	@R: overwriting insn
Data paths	ARF/ROB → RS RS → FU FU → ROB ROB → ARF	PRF → FU FU → PRF
Precise state	Simple: clear everything	Complex: serial/checkpoint

- R10K-style became popular in late 90's, early 00's
  - E.g., MIPS R10K (duh), DEC Alpha 21264, Intel Pentium4
- P6-style is making a comeback
  - Why? Frequency (power) is on the retreat, simplicity is important

# Speculation Recovery

- Squashing instructions in front-end pipeline



# Stall and Drain (1/2)

- Squash in-order front-end (as before)
- Stall dispatch (no new instructions → ROB, RS)
- Let OoO engine execute as usual
- Let commit operate as usual except:
  - Check for the mispredicted branch
    - Cannot commit any instructions after it
      - Any insns. in pipeline are on the wrong path
    - Flush the OoO engine
    - Allow dispatch to continue

# Stall and Drain (2/2)

- Delays recovery until BR retires

Ideal:

LOAD
ADD
BR
<i>junk</i>

LOAD
ADD
BR
XOR
LOAD
SUB
ST
BR

Stall & Drain:

LOAD
ADD
BR
<i>junk</i>

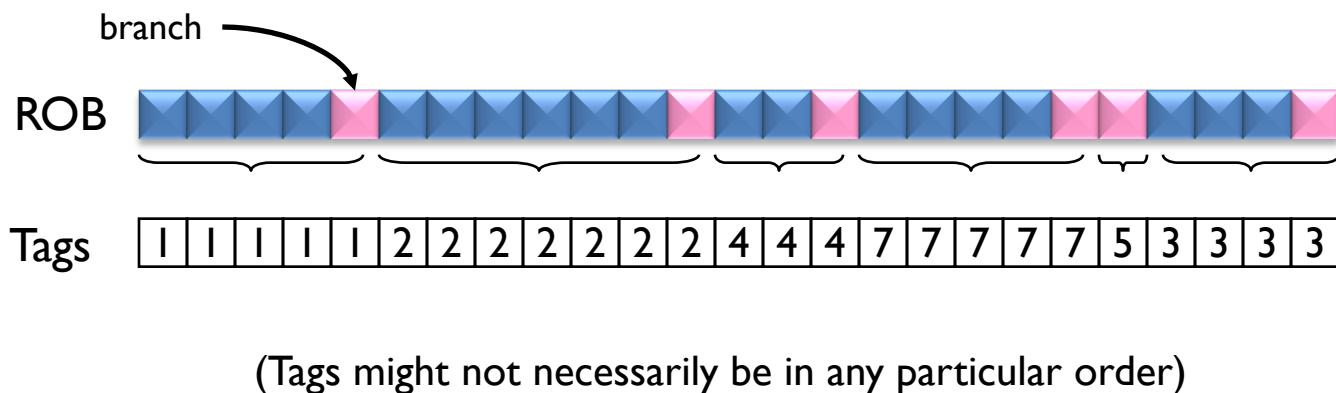
---
---
---
---
---
<i>junk</i>

---
---
---
---
---
XOR
LOAD
SUB
ST
BR

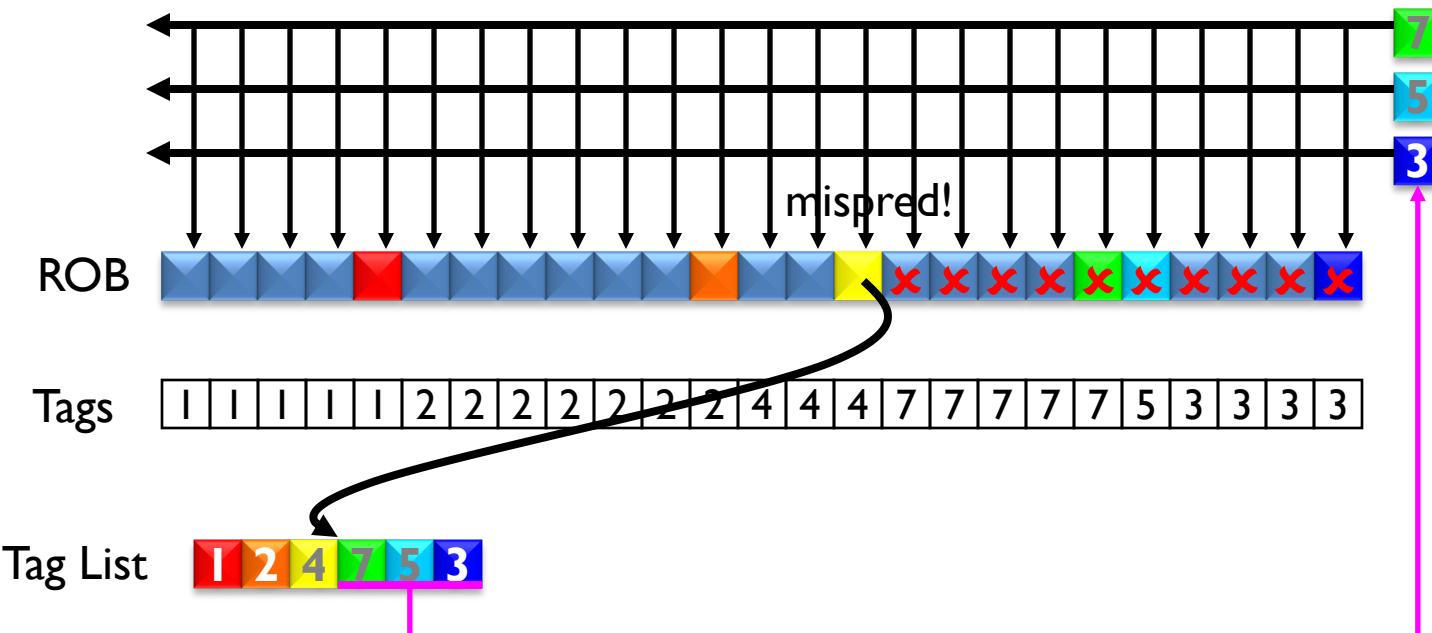
Simple to build, but low performance

# Branch Tags/Colors (1/2)

- Each insn. assigned the “current” branch tag
- Each predicted branch allocates a new branch tag
  - Newly allocated tag becomes current



## Branch Tags/Colors (2/2)



# Branch Tags for RS, overkill for ROB

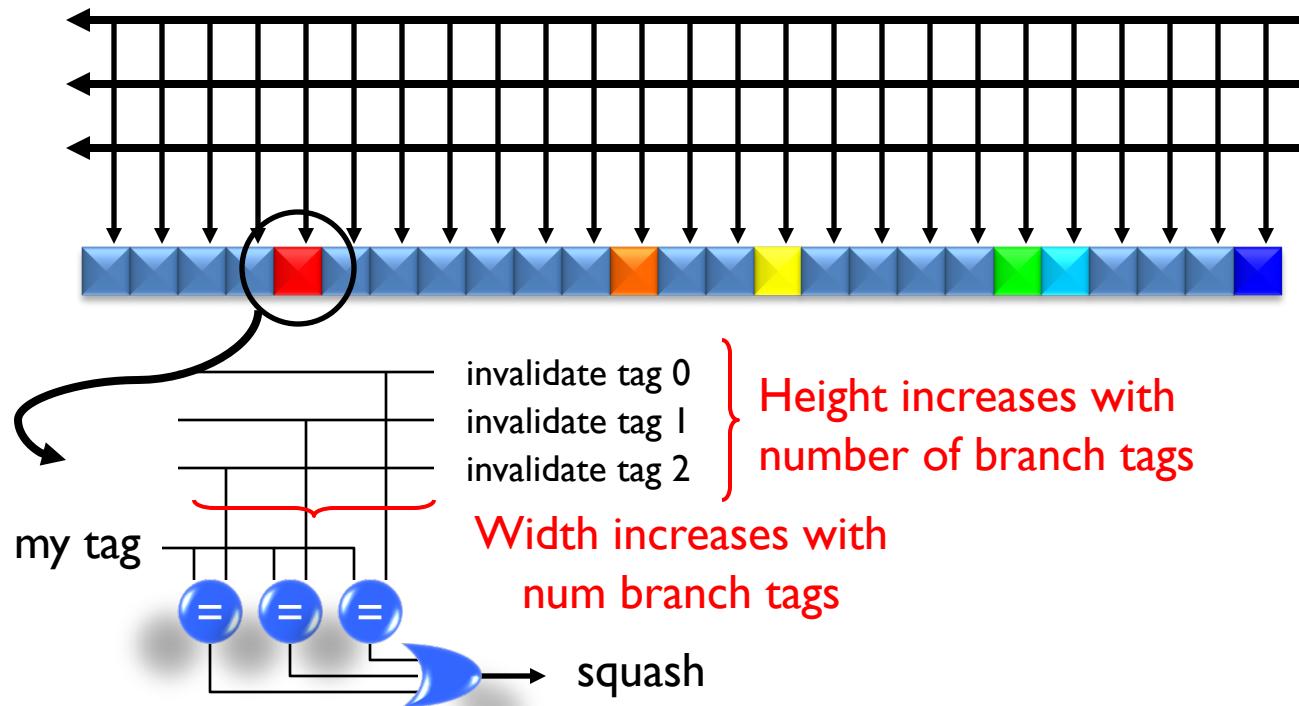
- ROB keeps insns. in program order
- Squash all insns. *after* mispredicted branch



- Tagging/coloring useful for RS
  - Insns. in RS are in arbitrary order
  - May be organized into multiple sets of RSs



# Hardware Complexity



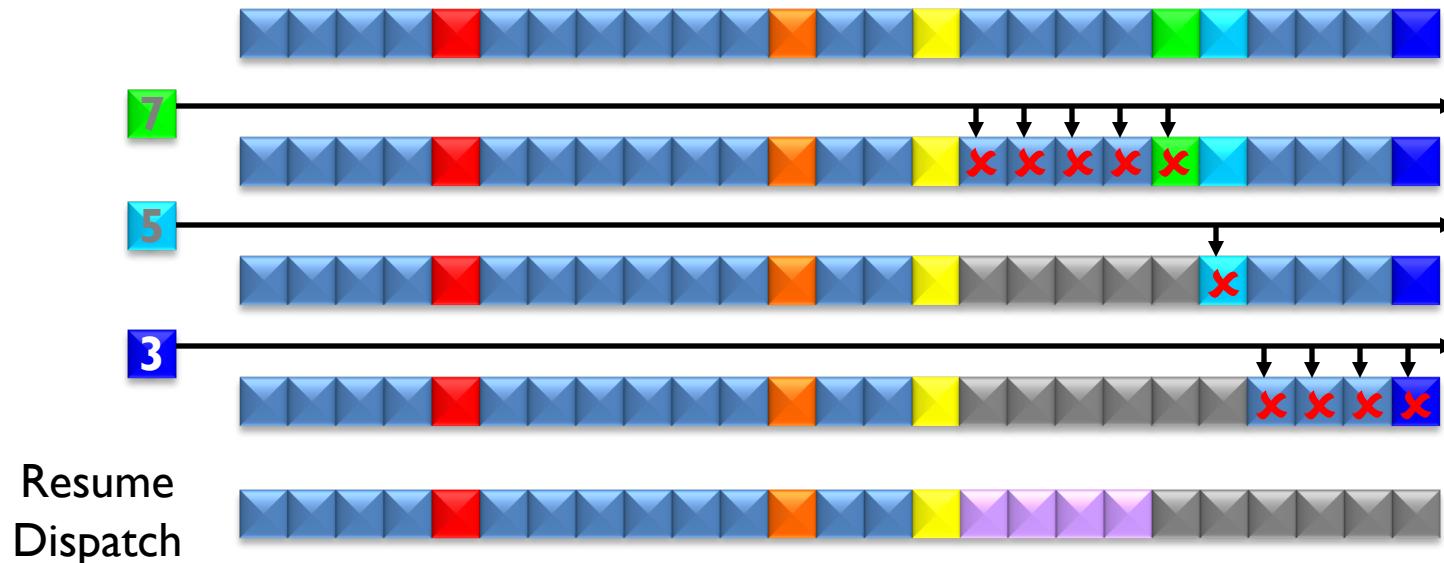
Area overhead is *quadratic* in tag count

# Squash Simplifications (1/2)

- For  $n$ -entry ROB, could have  $n$  different branches
- In practice, only a fraction of instructions are branches
  - Limit to  $k < n$  tags instead
  - If  $k+1^{st}$  branch is fetched, stall dispatch (structural hazard)

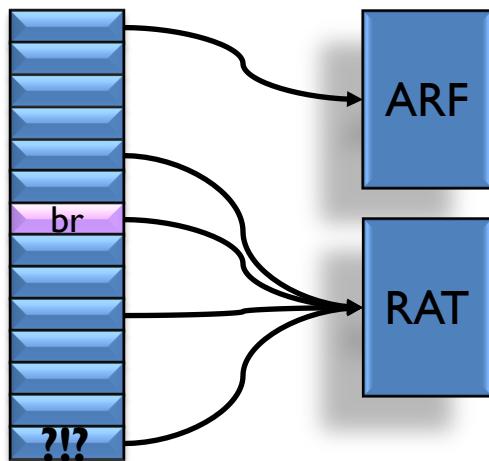
# Squash Simplifications (2/2)

- For  $k$  tags, need to broadcast all younger tags
  - Results in  $O(k^2)$  overhead
- Limit to few (e.g., one) broadcast per cycle



Can fetch and decode while squashing in back-end

# Register Speculation Recovery



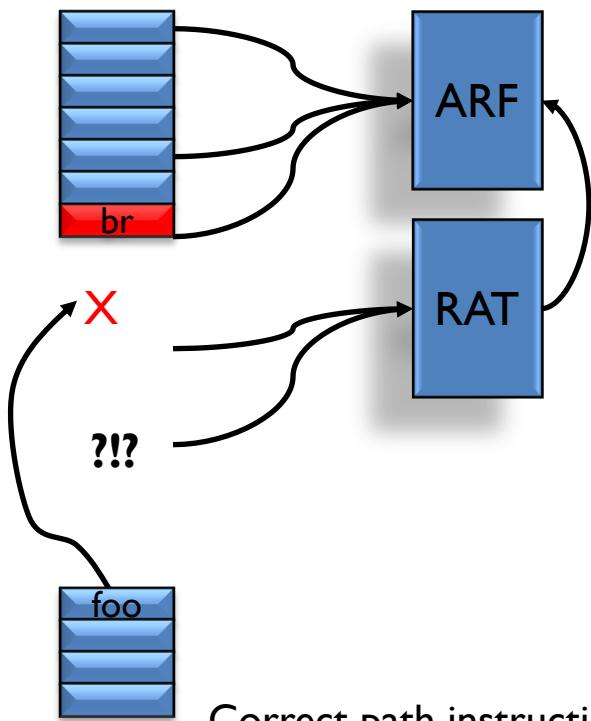
ARF state corresponds to state prior to oldest non-committed instruction

As instructions are processed, the RAT corresponds to the register mapping after the most recently renamed instruction

On a branch misprediction, wrong-path instructions are flushed from the machine

RAT left in invalid state

# Solution 1: Stall and Drain



Allow all instructions to execute and commit; ARF corresponds to last committed instruction

ARF now corresponds to the state right before the next instruction to be renamed (foo)

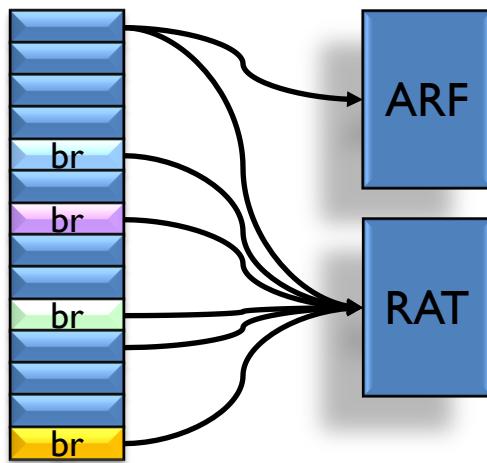
Reset RAT so that all mappings refer to the ARF

Resume renaming the new correct-path instructions from fetch

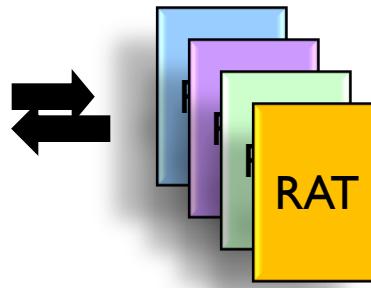
Correct path instructions from fetch;  
can't rename because RAT is wrong

Simple to build, but low performance

# Solution 2: Checkpointing (1/2)



At each branch, make a copy of the RAT  
(register mapping at the time of the branch)



Checkpoint  
Free Pool



On a misprediction:

1. flush wrong-path instructions
2. deallocate RAT checkpoints
3. recover RAT from checkpoint
4. resume renaming



Squash tags/colors can be same as checkpoints

## Solution 2: Checkpointing (2/2)

- No need to stall front-end
  - Need to “flash copy” RAT
    - Both for making checkpoints and recovering
  - Need to recover wrong-path checkpoints
- More hardware
  - Need one checkpoint per branch
  - What if the code has all branches?
    - Stall front-end when out of branch colors/checkpoints

# Solution 3: Undo List

- Each ROB entry tracks two physical registers
  - Its destination register
  - The previous physical register mapping
    - Required for R10K-style OoO anyway
- Walk ROB backwards, applying the old mappings
  - Low overhead: don't need full copies of the RAT
  - Slower: need to walk the ROB
  - Flexibility: can recover to any instruction
- Can combine with checkpointing
  - Checkpoint low-confidence branches; walk ROB for others